Conformity and Silence:

Experimental Evidence on Social Pressure and Free Speech*

Juan S. Morales† Margaret Samahita‡

July 31, 2026

Abstract

This paper studies the public expression of opinions on polarized topics (race and gender). In a series of pre-registered online experiments conducted in the United States in 2021 and 2023, we elicited participants' views on these topics and their willingness to publish their views online. We find an asymmetric U-shaped relationship between ideological stance and willingness to share views publicly: moderates were least willing to speak up, willingness rose toward both ideological extremes and was highest among respondents with progressive views. Participants also perceived publicly expressed opinions as more left-leaning than underlying private opinions. A priming treatment highlighting the risk of negative social media backlash had modest and statistically insignificant effects. A social-information treatment informing respondents that many others were willing to speak up significantly reduced self-censorship.

JEL Codes: D70, D83, P00, C90, Z13
Keywords: social media; spiral of silence; public opinion; cancel culture; free speech

*This paper was previously titled "Can Social Pressure Stifle Free Speech?". We thank the editor Daniele Nosenzo and two anonymous reviewers, Martin Bisgaard, Luca Braghieri, Leonardo Bursztyn, Eric Dickson, Edoardo Grillo, Caroline Le Pennec, Jean-Robert Tyran, Wieland Müller, Chris Roth, seminar and conference participants at the Burgundy School of Business, Collegio Carlo Alberto, Durham University, London School of Economics, Lund University, Trinity College Dublin, University College Dublin, University of Nottingham, University of Stirling, University of Vienna, Wilfrid Laurier University, Workshop on "Online Social Influence" (UCD), Workshop on "Directions of Polarization, Social Norms, and Trust in Societies" (MIT), 2022 APSA Meeting, 2022 ESA Special Meeting, 2022 SSES Annual Congress (Fribourg), 2022 ESA European Meeting (Bologna), 2023 Rebecca B. Morton Conference on Experimental Political Science (NYU), 2023 IMEBESS Meeting (Lisbon), 2023 EEA Meeting (Barcelona), and the 2023 CPEG Meeting (Queen's University), for productive discussions and helpful suggestions. The experiments described were approved by the ethics committees at UCD (HS-E-21-43-Samahita, HS-22-50-Samahita) and WLU (8354), and pre-registered as AEARCTR-0007905 (Samahita, 2021). Funding from UCD, the Collegio Carlo Alberto, LCERPA (WLU) and the Institute for Humane Studies (grant no. IHS017631) is gratefully acknowledged. All errors are our own.

†Department of Economics, Lazaridis School of Business and Economics, Wilfrid Laurier University. E-mail: jmorales@wlu.ca.

‡School of Economics and Geary Institute for Public Policy, University College Dublin. E-mail: margaret.samahita@ucd.ie.

1 Introduction

We often express public opinions which differ from our private views (Kuran, 1997). Other times, we may prefer to refrain from expressing our views altogether (Noelle-Neumann, 1974). Social norms, image concerns, and stigma are all factors which affect the public expression of opinions. With increased political polarization and the rise of social media, many prominent voices have recently argued that these social norms have become too strict and that fear of social backlash has led to a stifling of freedom of expression. In recent surveys, sixty-two percent of Americans agreed that "the political climate prevented them from saying things they believe because others might find them offensive" (Cato), and fifty-seven percent of UK residents reported that sometimes they stop themselves from "expressing their views on political and/or social issues because of fear of judgement or negative responses from others" (YouGov). These sentiments were echoed by a letter co-signed by numerous public figures, which argued that "the free exchange of information and ideas, the lifeblood of a liberal society, is daily becoming more constricted" (Harpers). Understanding when and why individuals choose to self-censor or conform is vital to diversity of thought and strong democratic institutions (Benson, 2024).

This paper reports results from a series of experiments in which participants report their views on controversial topics and their willingness to publish these views online. We first outline a simple model that formalizes the setting and motivates the empirical analysis. In the model, individuals trade-off espousing public views which are consistent with their true private opinions, while facing the cost of social stigma if these views differ from a socially accepted *expressive norm*. Higher perceived social costs of expression affect public expression in two ways: by shifting stated views toward the norm (*conformity*), and by reducing willingness to speak publicly (*silence*). The model also allows others' willingness to speak up to affect the utility of public expression, either by providing information about the strictness of the norm or by changing the perceived benefits of speaking alongside others with similar views.

We then conduct two pre-registered online experiments with US participants (total N=3,152) in which we elicit their views on divisive topics: support for the participation of trans women in competitive sports ("gender"), and opinions about whether other people are too sensitive regarding issues of race ("race"). These are highly polarized issues in the US and frequent sources of political conflict. We measure respondents' willingness to speak up based on their agreement to share their stated views on social media. Specifically, we ask participants whether they would be willing to let us post their opinion, along with their name, on a public Twitter page dedicated to the experiment. We also elicit participants' beliefs about the opinions of all other participants, as well as their beliefs about the opinions of participants who were willing to publish their views. These questions are incentivized.

We find an asymmetric U-shaped relationship between ideology and willingness to publicly post views: moderates are least willing to speak, while willingness increases toward both ideological extremes and is highest among participants holding opinions on the left. In addition, participants perceived the overall distribution of opinions to be more right-leaning than the distribution of publicly expressed views, implying that progressive positions were seen as less socially costly to voice. Speaking up is also significantly correlated with how much participants report self-censoring due to the political climate, and with their views on how important these topics are. Taken together, these descriptive patterns document systematic differences between private opinion, willingness to publish, and perceived public expression in a polarized US political environment.

Our experimental interventions provide further evidence on whether public awareness, perceived backlash, and social information can affect stated opinions and willingness to publish those opinions online. In a first intervention, participants were made *aware*, while reporting their views, that they might later be asked whether they are willing to have those views posted on social media (**Awareness** treatment). This intervention allows us to test whether the prospect of public disclosure affected participants' stated views, including whether it shifted responses toward views perceived as more socially acceptable. Publication awareness did not shift participants' reported views overall. However, among participants who did not receive a cancel-culture prime (described below), publication awareness shifted responses to the left (more progressive).

In our second treatment, participants were exposed to a *priming* text about "cancel culture", which described examples of individuals who lost their jobs following negative backlash over something they had posted on social media (**Prime** treatment). While we observe a negative effect of this intervention in the final wave of the experiment—participants becoming less willing to publish their views—priming participants to consider cancel culture and online backlash did not significantly change willingness to speak up overall.

Finally, in a third intervention, we informed participants about others' willingness to publicly express their opinions (**High/Low Peer Participation** treatments). In particular, we informed participants that for a previous group of participants "X% of them were willing to publish their opinion on the above statement with their name". We experimentally varied X to be high or low. Our results reveal increased willingness to speak up when a high share of others are also speaking up. This effect, however, was again concentrated among participants who had not been exposed to the backlash prime.

At the end of the experiment, we asked participants in an open text box about their reasons for choosing whether to speak up or not. Using a large language model (LLM) we analyze this text and classify these rationales into substantive categories. This analysis provides additional descriptive evidence on participants' stated reasons for publishing or not publishing their views. Participants with progressive views were less likely to cite fear of social backlash as a reason not to post. Moderates more often described themselves as uninterested in the issues or disengaged from social media, while conservatives were more likely to cite free speech as a reason to speak up. These patterns are consistent with the interpretation that progressive views were perceived as facing lower expressive costs.

Our paper makes several important contributions. Unlike prior studies focusing on university or academic contexts (Norris, 2023; Braghieri, 2024; Ho and Huang, 2024; Braghieri, Bursztyn and Fasnacht, 2026), we provide evidence from a broader online sample representative of the US adult population, contributing new evidence to debates surrounding cancel culture. In addition, our finding that moderates were least willing to publish their views online is consistent with research showing that extreme views are overrepresented on social media (Robertson, Del Rosario and Van Bavel, 2024). We also document how willingness to speak up significantly increases when individuals learn that a large proportion of previous participants are also publicly expressing their views, highlighting the potential role of social cues in overcoming self-censorship.

Our work relates to the large body of literature on social image and reputational concerns (Bernheim, 1994; Bénabou and Tirole, 2006; Andreoni and Bernheim, 2009). Individuals often misreport their true opinion (Kuran, 1997; Braghieri, 2024; Valentim, 2024) or change their behavior (Andreoni and Petrie, 2004; Rege and Telle, 2004; Friedrichsen and Engelmann, 2018; Schulz, 2025) in the presence of an audience due to social image concerns and a desire to conform to social norms. Fear of social stigma may also result in individuals self-censoring their opinion (Noelle-Neumann, 1974; Morales, 2020; Gibson and Sutherland, 2023; Bursztyn et al., 2023) or hiding behaviors that deviate from the perceived social norm (DellaVigna et al., 2016; Carlson and Settle, 2016; Holm and Samahita, 2018; von Siemens, 2020). However, the specific role of cancel culture in shaping expressive norms and behavior remains relatively under-explored in the existing literature.

Social image concerns are particularly relevant in online settings, where backlash and moral outrage, often linked to cancel culture, can intensify these pressures (Brady et al., 2021; Crockett, 2017; Forestal, 2024). Recent work has documented how perceptions of cancel culture are highest for academics who hold minority opinions (Norris, 2023, 2025) but are generally exaggerated in the wider population (Dias, Druckman and Levendusky, 2025). These issues also relate to the broad debate regarding political correctness and the value of free speech (Morris, 2001; Voerman-Tam, Grimes and Watson, 2023). Our priming intervention, which exposes participants to consider cancel culture and the potential for online social backlash, allows us to study the extent to which this particular form of social pressure affects their willingness to speak up. We show that increased awareness of negative backlash on social media had only moderate effects on self-censorship.

Our work also contributes to the literature on political participation, particularly how social pressure and social cues influence individuals' willingness to engage in public discourse and activism. Social cues can affect reported attitudes, political donations, voting, and other behaviors (Bond et al., 2012; Perez-Truglia and Cruces, 2017; Bursztyn and Jensen, 2017; Conzo ⓡ et al., 2023; Isler and Gächter, 2022). In the context of dissent, observing others speak up can motivate individuals to voice their own opinions and participate in activism (González, 2020; Manacorda and Tesei, 2020), thereby eroding strict expressive norms (Morales, 2020; Bursztyn, Egorov and Fiorin, 2020; Álvarez-Benjumea, 2023; Albornoz, Bradley and Sonderegger, 2022; Dinas, Martínez and Valentim, 2024; Apffelstaedt, Freundt and Oslislo, 2022). Relatedly, Ho and Huang (2024) show that highlighting the presence of self-censoring individuals leads to higher levels of speaking up among those who hold dissenting views. At the same time, others' participation can also lead to free-riding and act as a strategic substitute for individual dissent (Cantoni et al., 2019; Hager et al., 2023). Our study shows that individuals are more likely to speak up when *many* others do so, even without knowing what views are being expressed. These findings complement evidence that political polarization distorts social behavior through norm-based mechanisms (Dimant, 2024), by showing that information about others' willingness to publish can increase participants' own willingness speak up.

In the following section we outline a simple conceptual framework that links perceived social costs to stated views and willingness to publish. Section 3 explains the experimental design and our implementation. Section 4 details the results of our experiments. Section 5 documents empirical extensions and Section 6 concludes.

2 Conceptual Framework

2.1 Model

To motivate our hypotheses and fix ideas for our experimental interventions, we present a stylized model in which individuals trade-off voicing their private views against a social norm n which regulates public expression. In the model, this *expressive* norm is defined as the position that can be expressed at the lowest social cost. The framework captures and formalizes two potential effects of changes in expressive costs on public speech: conformity and silence. Other recent theoretical explorations of related issues include (i) studies of *conformity*, such as the effect of changing norms on public expression (Bursztyn, Egorov and Fiorin, 2020), the role of social image in driving conformity (Braghieri, 2024), and the emergence of biased norms under peer pressure (Michaeli and Spiro, 2017), and (ii) studies of *silence*, such as the salience of silence and its influence on public expression (Ho and Huang, 2024), and the effect of justification on dissent (Bursztyn et al., 2023). Our model captures the interplay between expressive costs, conformity, and silence in a simple unified framework.

Assume a continuum of individuals N defined over their private opinions $o_i \in [0,1]$. Individuals consider whether they want to publicly express an opinion on a particular issue. Specifically, individuals choose: i) a public stance $s_i \in [0,1]$, and ii) whether to "speak-up" and voice their stance, $v_i \in \{0,1\}$, given that there is an expressive norm that dictates what an appropriate public stance is: $n \in [0,1]$.

In reality, expressive norms are shaped by the social environment, including who speaks up—in that sense, n is ultimately an equilibrium object. However, for exposition purposes we keep a common n and for simplicity we treat it as exogenous, which we view as a useful way to think about the behavior of our survey participants, who take these norms as given. This approach also follows a broad literature that treats social norms, identity prescriptions, and perceptions of appropriate behavior as exogenous constraints shaping individual behavior (e.g. Akerlof and Kranton, 2000; Kimbrough and Vostroknutov, 2016; Bašić and Verrina, 2024; Hauk and Ortega, 2025; Acemoglu, Egorov and Sonin, 2026).[1]

Individuals also get an intrinsic payoff for speaking up plus some idiosyncratic component: $\kappa + \varepsilon_i$, where $\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$. They maximize their utility:

$$u_i = \begin{cases} -[\alpha(s_i - o_i)^2 + \beta(s_i - n)^2] + \kappa + \varepsilon_i & \text{if } v_i = 1 \\ 0 & \text{if } v_i = 0 \end{cases} \tag{1}$$

where β is the cost or risk from social disapproval and α is a "cognitive dissonance" cost from expressing a public stance which differs from the individual's private opinion.[2]

The maximization problem can be solved by backward induction, first choosing the optimal public stance s_i^*, and then whether to speak up or not v_i^*. At the unique optimum,

[1]Norms may also vary by individual if, for instance, Democrats and Republicans adhere to different norms. However, allowing for heterogeneous norms does not change our main predicted treatment effects.

[2]In some contexts, silence itself may be an informative signal about individual views (Ho and Huang, 2024), but we abstract from this possibility because it is not relevant in our experimental context.

individuals choose:

$$s_i^* = \frac{\beta n + \alpha o_i}{\beta + \alpha}$$

The costs of speaking up are minimized when there is no conflict between the private opinion and the expressive norm, $s_i^* = o_i = n$. When $o_i \neq n$, s_i^* will be a weighted average between the private opinion and expressive norm, with weights that depend on the social disapproval and cognitive dissonance parameters β and α.

The utility of expressing the above optimal public stance is

$$u_i(v_i = 1) = \kappa + \varepsilon_i - \frac{\beta\alpha(o_i - n)^2}{\beta + \alpha} \tag{2}$$

Thus, the individual will speak up if

$$(o_i - n)^2 \leq \frac{\beta + \alpha}{\beta\alpha}(\kappa + \varepsilon_i) \tag{3}$$

Define $v^*(o)$ as the share of individuals who find it optimal to speak up out of those holding private opinion o. Assuming that the idiosyncratic payoff component ε is independent of private opinion o, the private opinion associated with the highest willingness to speak up coincides with the expressive norm n:

$$\arg\max_o v^*(o) = n \tag{4}$$

Equation 4 should be interpreted as an identification result rather than a comparative-static prediction: because the expressive norm n is defined as the view that can be expressed at the lowest social cost, the model implies that individuals whose private opinions are closest to n are most willing to speak publicly.[3]

Our comparative statics study changes in β, an increase in which corresponds to stricter norms regarding public expression. We consider two main outcomes: individuals' reported views s and their willingness to publicly express those views v—which match our experimental setting. In turn, we refer to changes in these actions respectively as increased *conformity* (s) and *silence* (v). Based on this simple framework, we highlight three insights related to the public expression of opinion and social norms that regulate speech.

[3]See Appendix A.1 for proof and an extension allowing expressive utility from ideological extremeness, for example because these views are more identity-relevant or perceived as especially important (Golman, 2023; Le Yaouanq, Schwardmann and van der Weele, 2025). Under this extension, the assumption of ε being independent of private opinion o may be violated, so the willingness-to-speak peak can identify the expressive norm only under additional restrictions, as discussed in the appendix.

Observation 1 *(Conformity):* As expressive norms become more strict (β increases), conformity increases. That is, an increase in β reduces the distance between expressed views s and the norm n.

To see this, note that the absolute distance between the optimal public stance (D) and the norm is:

$$D_i^* = |s_i^* - n| = \frac{\alpha}{\beta + \alpha}|o_i - n|$$

The effect of increasing β on individuals' public stance s_i^* is thus to move it closer to the norm n (that is, $\partial D_i^* / \partial \beta < 0$).

Observation 2 *(Silence):* As expressive norms become more strict, self-censorship increases. That is, an increase in β reduces the likelihood that individuals choose to voice their views publicly (v).

If individuals choose to speak up, that is, if $v_i = 1$, then:

$$u_i(v_i = 1) = -[\alpha(s_i - o_i)^2 + \beta(s_i - n)^2] + \kappa + \varepsilon_i$$

and this utility decreases as β increases (that is, $\partial u_i(v_i = 1)/\partial \beta < 0$). As the payoff of speaking up decreases, it increases the likelihood that individuals cross the threshold ($u_i(v_i = 1) < u_i(v_i = 0)$) at which they would rather remain silent.

Observation 3 *(Differential silence):* The silencing effect of stricter expressive norms (as represented by an increase in β) is more pronounced for individuals whose public stance s_i deviates further from the norm n.

Formally, the sensitivity of utility to a change in β increases as the absolute difference between s_i and n increases. Note that:

$$\frac{\partial u_i}{\partial \beta} = -(s_i - n)^2 \tag{5}$$

That is, the effect of an increase in β on reducing the likelihood of speaking up is stronger for individuals whose public stance s_i deviates more from the norm n.[4] Note also that stricter norms (an increase in β) do not affect individuals whose private opinion o_i matches the expressive norm n, since this opinion is not stigmatized (and they would choose $s_i^* = n$).

[4]The cross partial derivative of the utility function with respect to β and $|s_i - n|$ is negative. In particular, we can consider two cases: (1) If $s_i > n$, the impact of an increase in β on the decision to speak up becomes more negative as s_i increases: $\frac{\partial^2 u_i}{\partial \beta \partial s_i} = -2(s_i - n) < 0$. (2) If $s_i < n$, the impact of an increase in β on the decision to speak up becomes more negative as s_i decreases: $\frac{\partial^2 u_i}{\partial \beta \partial s_i} = -2(n - s_i) < 0$.

2.2 Social participation

Now consider a simple extension of the model in which individuals observe how many others are publicly expressing their views. We consider this a signal that allows agents to learn about the social disapproval costs (β), and that could affect the intrinsic reward for speaking up (κ).[5]

The utility of speaking up is now given by:

$$u_i = -[\alpha(s_i - o_i)^2 + \beta(v_{-i})(s_i - n)^2] + \kappa(v_{-i}) + \varepsilon_i$$

In this context, the social costs $\beta(v_{-i})$ are decreasing in v_{-i}, the proportion of others who speak up: when others publicly express their views, individuals can infer that social sanctioning is weak. However, the effect of peer participation v_{-i} on the reward for speaking up $\kappa(v_{-i})$ is ambiguous due to the possibility of both strategic complementarity or substitutability in voicing one's views, as in the literature on collective action (Cantoni et al., 2019; Hager et al., 2023). Thus, the overall effect of higher peer participation

$$\frac{\partial u_i}{\partial v_{-i}} = -\beta'(v_{-i})(s_i - n)^2 + \kappa'(v_{-i})$$

will depend on whether any strategic substitutability ($\kappa'(v_{-i}) < 0$) is sufficient to undo the positive effect of weaker perception of social sanctioning coming from the first term.

3 Experimental Design

3.1 Motivation and context

We study public expression amid online backlash and social image concerns during a period of prevalent discussions about cancel culture. The US provides a relevant setting given rising political polarization (Boxell, Gentzkow and Shapiro, 2024; Draca and Schwarz, 2024; Dimant, 2024) and an increase in self-censorship (Cato). Sensitive topics like race and gender norms amplify these pressures, especially on social media, where expressions of stigmatized opinion could carry social costs. To examine these issues, we conducted experiments in 2021 and 2023 (pre-analysis plans available here; timelines in Figures 2 and 3).

[5] In reality, observing others' public expression may also affect perceptions of the expressive norm. For simplicity, we leave this channel outside the formal extension, but we investigate this possibility empirically in Section 5.1.

3.2 Main variables

Our outcomes of interest are reported attitudes (which we refer to as s) and participants' willingness to share their views on social media (v).

3.2.1 Elicitation of Attitudes (s)

Participants were asked to consider two statements in random order:

- In my opinion, trans women should be allowed to participate in women's sports competitions.
- In my opinion, many people nowadays are too sensitive about things to do with race.

Responses range from *strongly disagree* to *strongly agree* (coded as 1-7).[6]

3.2.2 Elicitation of Willingness to Publish (v)

We asked two questions about sharing participants' responses on social media, in order:

1. **Posting anonymously:** "Would you be willing to let us post on social media, anonymously, your response to the previous statement:" (for example)

   [Participant 37]
   *"**I somewhat disagree**" that many people nowadays are too sensitive about things to do with race."*

   Participants were informed that selecting Yes meant we would create a tweet containing the above and post it on a public Twitter page after data collection.[7]

2. **Posting with own name:** "Would you be willing to let us post on social media, together with your name, your response to the previous statement:" (for example)

   [Your name here]
   *"**I somewhat disagree**" that many people nowadays are too sensitive about things to do with race."*

   Participants were informed that:

[6] Statements were pre-tested in a pilot together with other statements, see Appendix Figures A8 and A9. In Experiment 1, we also included a placebo statement—"People who have been vaccinated against COVID-19 should be allowed to travel without testing and quarantine requirement"—and found no treatment effect on this outcome.

[7] Some anonymous responses were posted at https://twitter.com/SurveyResponses, but our tweeting program was later blocked by Twitter due to spam-like behavior (understandably).

- We will create a tweet containing the above response and may post it on a public Twitter page created once data collection is complete (* see below)
- *We will contact Prolific to request your first and last names. Note that while in general Prolific does not allow researchers to collect personal information, Prolific does encourage researchers to get in touch in cases such as this, where the study design requires the collection of personal data (see https://researcher-help.prolific.co/hc/en-gb/articles/360015378834-Can-I-ask-Participants-for- their-Personal-Information-Identifiers-).[8]
- The tweet will only contain a text of your name without any hyperlink, the public Twitter page will potentially contain the names and opinions of many participants.
- The link to the public Twitter page will be made available to participants who contact the researcher to ask for it, but it will not be otherwise advertised. The public Twitter page will be deleted after 30 days.

Full instructions are in Appendix D.

Since we study the effect of perceived social cost on opinion expression, it was essential that participants seriously considered the possibility of posts being made public. However, we did not want to actually publish names due to potential negative consequences for participants. In our aim of aligning with norms of no deception we truthfully informed participants that if they consented we would attempt to obtain their names, but publication was conditional on an event with an extremely low probability, as we then explained in the debrief (see Section 3.5 below). Previous requests to Prolific for participant names have been denied, as expected.

This design allowed us to measure real behavioral responses while avoiding deception and minimizing harm. The experiments were approved by two separate ethics committees.

3.3 Experimental interventions

We implemented three interventions designed to vary social pressure, or the cost of "speaking up" against a prevailing expressive norm. The **Prime** treatment primed participants on the risks of online backlash. The **Publication Awareness** treatment informed them that their responses could be published, increasing the saliency of social backlash. The **High**

[8] At the time, Prolific did not explicitly forbid the collection of personal data, stating: "There may be some other cases where your study design requires the collection of personal data. In this case, please get in touch with Prolific Support to discuss which approaches might be possible." (Link, accessed 2025-05-29).

Peer Participation (HiPeer) treatment (and LoPeer as an active control) indicated whether a high/low percentage of previous participants agreed to publish. Below, we describe each treatment together with a brief discussion of our pre-registered predictions. Further details are available in the pre-registration.

3.3.1 Prime treatment

This intervention primed participants on the potential risk of social disapproval by displaying the following text and image:

The public nature of social media has resulted in individuals sometimes experiencing negative consequences as a result of their posts, in a phenomenon that some people refer to as **"cancel culture"**.

*"Those most vulnerable to harm tend to be **individuals previously unknown to the public**, like the communications director who was **fired** in 2013 after posting on social media, from her personal account, **an ill-thought-out joke** about Africa, AIDS and her own white privilege ... or the data analyst who was **fired** last spring after posting on social media, after the death of George Floyd in police custody, a study that suggested that riots depressed rather than increased Democratic Party votes."*

These cases highlight the risk of **public backlash from social media.**

Figure 1: Prime treatment

The text, adapted from a New York Times article on cancel culture, highlights the risks of expressing personal opinions online, which is expected to reduce participants' willingness to publish. Experiment 2 included a weaker variant (WeakPrime, Figure 3) to test whether any effect is driven by the term "cancel culture" or the accompanying image.[9]

[9]Details in Appendix Section A.2. The effects of Strong (original) and Weak primes are not statistically different, so we pool them in the analysis. While there may be concerns that these two priming interventions

We included an attention check, asking: "what does the text say cancel culture can result in?" Participants had to select "losing a job" before proceeding.

For the control text, participants saw information about University College Dublin (UCD) or the University of Turin (UniTo), with the university logo replacing the cancel culture image. Control texts were randomized, each with a corresponding attention check.

We also vary the timing of the prime, as shown in Figure 2. **Treatment 1** shows the prime *before* attitude elicitation. **Treatment 2** shows the prime *after* attitude elicitation (to investigate whether individuals who revealed their attitudes under no prime were more sensitive to publishing ex-post).

3.3.2 Publication Awareness treatment

In a second treatment dimension of Experiment 1, we varied participants' awareness of the potential publication of their views—revealing opinions to the public (T4-T5) instead of just the experimenter (T1-T3). Publication awareness creates an additional channel of social pressure, as participants perceive public exposure of their views as carrying higher social costs, which in turn is expected to reduce their willingness to publish.

In Treatments 1–3 (*Publication Unaware*), attitudes and willingness to publish were elicited in separate stages. In Treatments 4–5 (*Publication Aware*), both were elicited within the same stage (on separate pages), allowing participants to revise their stated opinion after indicating their willingness to publish it. Additionally, participants in T4–T5 were already aware of the possibility of publication when responding to the *second* opinion statement. This design allows us to test whether awareness of potential public exposure affects the content of expressed views.

To examine interaction effects, T4 included the online backlash prime before attitude elicitation, while T5 used a control text instead.

3.3.3 High/Low Peer Participation treatment

In a second intervention of Experiment 2, we varied the information given about how likely previous participants were to share their opinions. We interpret this manipulation as a reduced-form change in the perceived social environment, which may affect multiple factors influencing willingness to express. Higher peer participation may signal lower social sanctions, but it may also affect beliefs about others' behavior. In our experimental

could induce experimenter demand effect, it is unclear ex-ante which direction it would go and may depend on political affiliation. De Quidt, Haushofer and Roth (2018) show that typical demand effects are probably small. Similarly, Mummolo and Peterson (2019) find little evidence for demand effects in online survey experimental settings like ours.

design, we do not specify *what* views others are expressing, only the proportion of individuals willing to publish, but we elicit participants' beliefs about the views others express after the intervention. Additionally, in a collective action setting, expressing an opinion can be costly, leading to free-riding (strategic substitutes) or coordination benefits (strategic complements). The net effect is theoretically ambiguous (Cantoni et al., 2019; Hager et al., 2023).

Before asking about willingness to publish, participants saw:

> When we asked a previous group of participants, X% of them were willing to publish their opinion on the above statement with their name.

where X was randomized into either the "LoPeer" (13%/7% for gender/race) or "HiPeer" (60%/67% for gender/race) treatment. These values were drawn from previous surveys in US states with similar high/low willingness to publish, making LoPeer an active control for HiPeer. Participants were debriefed on this at the end of the experiment.[10]

3.4 Additional questions

Value of publishing: We also measured participants' willingness to pay to publish, with their name, their view on one randomly chosen statement (Gender or Race). Participants received 10 tickets for a USD 100 lottery (they were informed of their odds, which depended on the wave and number of participants). If they agreed to publish, they were asked if they would give up 10, 5, or 1 ticket to not post, sequentially. If they declined all, WTP was coded as 0. Those who initially refused to publish were offered 1, 5, 20, or 50 additional tickets to reconsider. If they still declined, WTP was coded as -100.

Topic importance and knowledge: Participants also rated the importance of each of the two issues (1 = Not important, 5 = Extremely important). In Experiment 2, they also rated their topic knowledge (1 = Little to no knowledge, 5 = Expert).

Beliefs about stated views: To measure participants' beliefs about the views expressed by other participants, and the extent to which publicly expressed views differed from private views, we asked about the average opinion of (i) all participants and (ii) those willing to publish without payment (for one randomly chosen statement). In Experiment 2, we instead elicited the *majority* opinion (Krupka and Weber, 2013). For instance, if participants perceive the overall opinion as more right-leaning than that of those willing to publish,

[10] LoPeer: [13%/7%] of [Maryland/Indiana] participants; HiPeer: [60%/67%] of [Arizona/Oregon] participants.

they perceive that left-wing views are more publicly accepted. Correct answers earned 5 extra lottery tickets.[11]

Concerns about political correctness: We also asked participants five questions regarding their views on the political climate regarding online freedom of speech and social pressure. For example: "The political climate these days prevents me from saying things I believe because others might find them offensive." (7-point scale, Strongly disagree - Strongly agree) and "How often do you think social pressure causes people to refrain or abstain from expressing political opinions on social media?" (7-point scale, Never - Always).

Demographic questionnaire: In both experiments, we collected data on demographics, risk attitude, political preferences, and social media use, eliciting the latter two pre-treatment due to potential heterogeneity (Montgomery, Nyhan and Torres, 2018). Political preference was measured via self-placement on a left-right scale and party affiliation, while social media use was assessed by daily time spent on various platforms Facebook, Twitter, Instagram and other social media platforms.

In Experiment 2, we added the Hong psychological reactance scale (Hong and Faedda, 1996), how common participants think their first and last name are (7-point scale), and the likelihood of participants' opinions truly being posted on social media (7-point scale).[12] Participants also provided open-ended reasons for deciding whether to publishing or not, and could give general feedback.

3.5 Debrief

We conclude by debriefing participants on the purpose of the study and informing them that any opinions approved for anonymous publication would be posted on a public Twitter page. Regarding publication with names, we inform participants that: "Previous requests to Prolific asking for participant's names in a similar study design have been turned down; so we do not anticipate that we will publish your opinion with your name, even if you stated that you would like us to do this. Regardless, if you stated that you were willing to publish the opinion with your name in exchange for lottery tickets, you will still get these additional lottery tickets." The full survey is in Appendix D.

[11] Results for average and majority opinion are similar (Table 2). Participants were unaware of alternative treatment arms.

[12] Only 13% of participants believed that it was "extremely unlikely" that responses would end up being posted. Results are not qualitatively different when excluding these participants, see Appendix Tables A12 and A15.

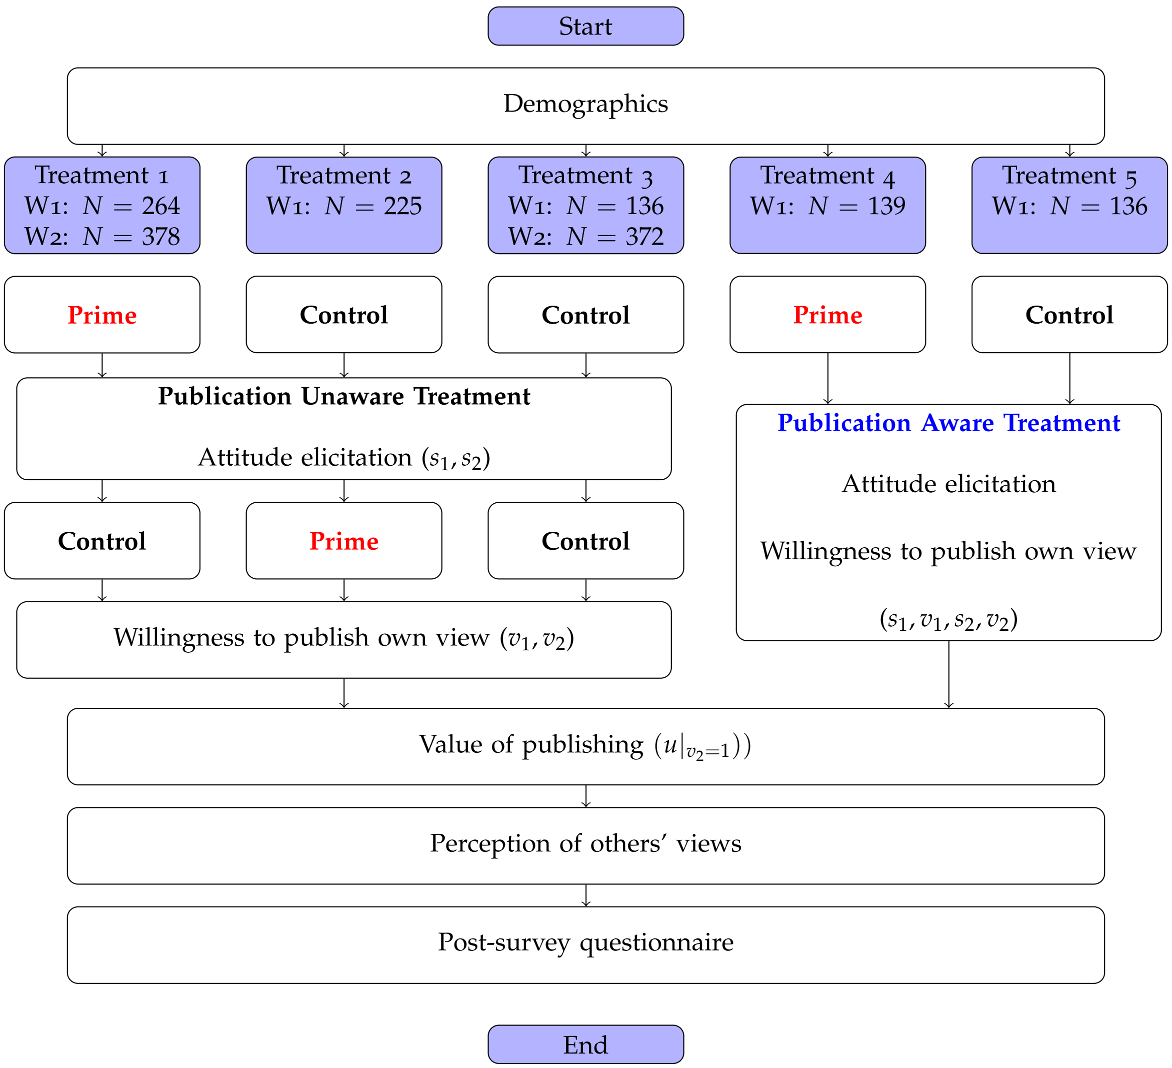

Figure 2: Experiment 1 timeline

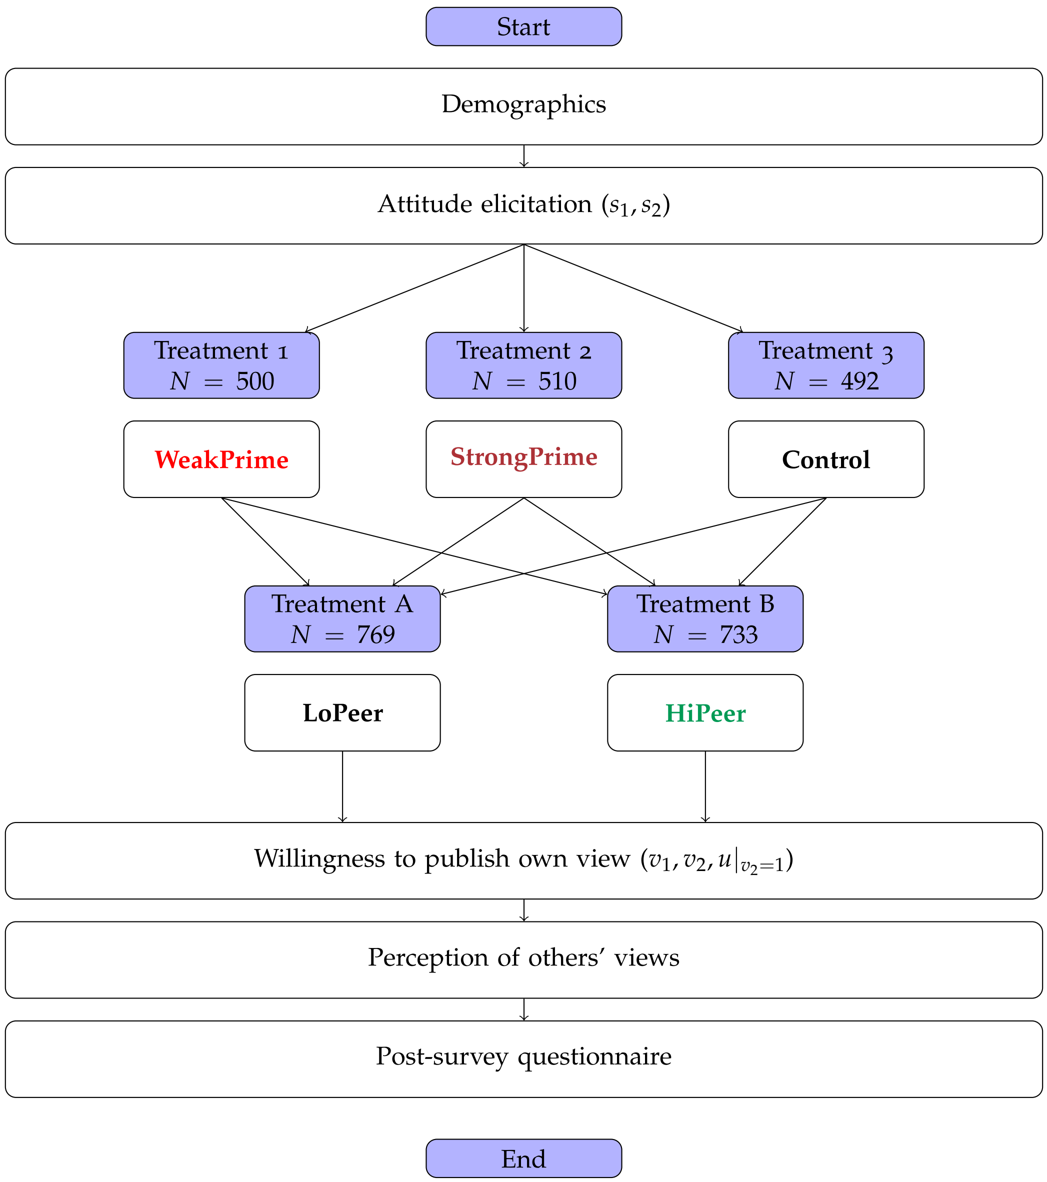

Figure 3: Experiment 2 timeline

3.6 Implementation

For Experiment 1, the first wave of data collection (Wave 1) took place in August 2021 via Prolific, recruiting 900 participants: 300 each from self-identified Democrats, Republicans, and Independents (including those with affiliation "Other" or "None").[13]

Wave 1 unintentionally over-sampled young females due to Prolific-specific sampling issues.[14] To address this, Wave 2 (Nov 2021) recruited a US nationally representative sample in terms of age, gender and ethnicity (750 participants), with the following modifications: (i) running only Treatments 1 (primed) and 3 (control) with equal numbers due to budget constraints, (ii) using only the gender question as it resulted in more distinct norms across partisan groups, and (iii) adding the question on perceived likelihood of actual publication.

For Experiment 2, to study the effect of peer participation, Wave 3 (March 2023) recruited 1502 participants: 400 Democrats, 700 Independents (including unaligned), and 400 Republicans—oversampling Independents to examine a potential backlash effect found in Experiment 1 (but which was not confirmed by the second experiment). Participants were randomized evenly into a 3x2 design (StrongPrime/WeakPrime/Control x HiPeer/LoPeer).[15]

Key descriptive statistics are in Appendix Table A1. As stated above, Wave 1 participants skew younger, more female, and more white than the nationally representative Wave 2. They are also more active on social media (80% spending ≥ 1 hour daily vs 55% in Wave 2). Waves 3 and 4 more closely resemble Wave 2 demographically, despite not being nationally representative.

4 Results

4.1 Descriptive evidence

Perceptions of political climate

Perceptions of the political climate around free speech reflected broader public narratives prevalent in popular media at the time, with many respondents reporting concerns con-

[13] We allocated 27.5% of participants to T1-2 and 15% to T3-5 to maximize power for our main tests. Using one-sided tests with 5% significance, we had 80% power to detect and effect size $\geq$ 0.20sd for our main test of H1 and and effect size $\geq$ 0.24sd for H2.

[14] https://blog.prolific.co/we-recently-went-viral-on-tiktok-heres-what-we-learned/, accessed 2021-10-13.

[15] To investigate whether the above absence of backlash effect was due to the addition of the High/Low Peer Participation statement, we conducted Wave 4 as a robustness check. Wave 4 replicated Wave 3 with an additional control treatment where participants were uninformed about others' participation (as in Experiment 1), using Independent participants. The results confirmed those found in Wave 3.

Figure 4: Perceptions of political climate and freedom of speech online

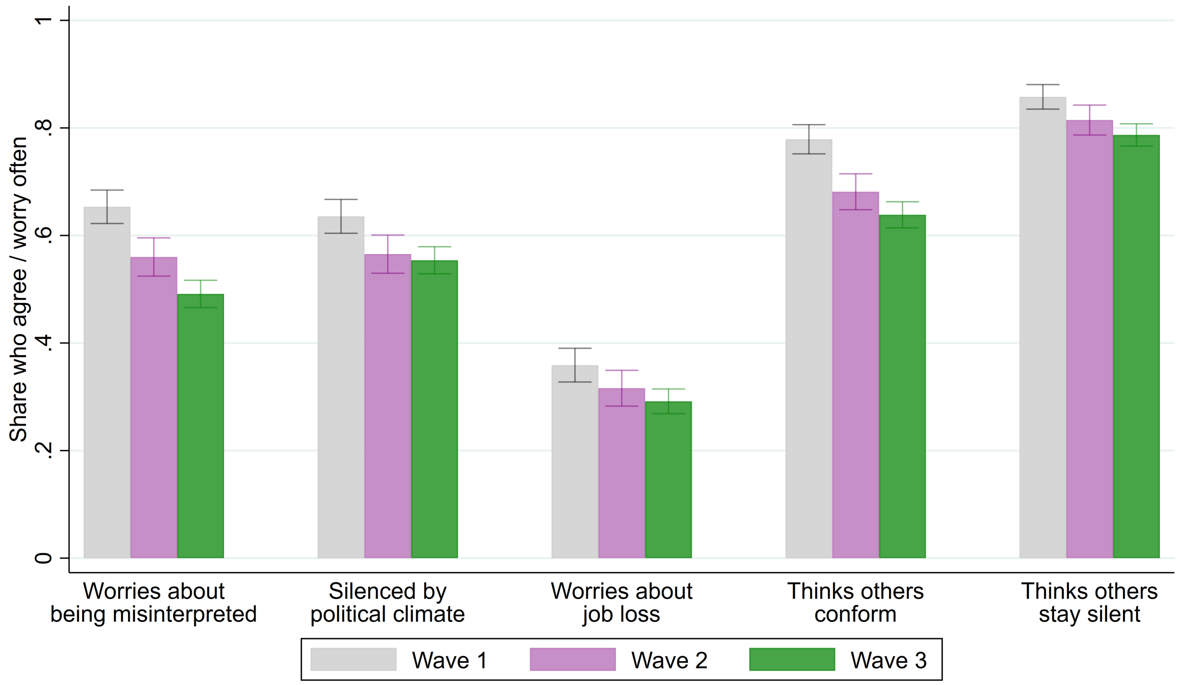

Notes: Respectively, the questions correspond to: "How often do you worry that things you post on social media can be misinterpreted?"; "The political climate these days prevents me from saying things I believe because others might find them offensive."; "Are you worried about losing your job or missing out on job opportunities if your political opinions become known?"; "How often do you think social pressure causes people to misrepresent or lie about their political opinions on social media?"; "How often do you think social pressure causes people to refrain or abstain from expressing political opinions on social media?".

sistent with discussions of "cancel culture" and censorship. Most participants believed that social pressure led individuals either to misrepresent their political views or to remain silent on social media: around eighty percent of respondents think that social pressure causes people to abstain from expressing their political views. Similarly, many report worrying about being misinterpreted online and agree that the current political climate constrains their ability to express their beliefs. Although fewer respondents report fearing job-related repercussions for their political opinions, a substantial minority nonetheless expressed such concerns (Figure 4).[16]

[16]In aditional analyses, we construct an index of concern for freedom of speech using the first principal component of the responses to the five questions on political correctness. Concern is higher among participants stating conservative views, active social media users, college graduates, and those more willing to take risks (Appendix Table A2). However, this index does not significantly vary across our experimental interventions (Appendix Tables A3 and A4).

Stated attitudes and willingness to speak-up

We next present descriptive evidence on attitudes toward gender and race topics and individuals' willingness to share their views online. Figure 5 shows the distribution of attitudes (s) on a 1-7 scale, where the gender topic was reverse-coded so that 1 represents more progressive views, split by willingness to publish opinions with names. These topics are very divisive, as revealed by a bimodal distribution of reported views.[17]

Figure 5: Public expression and attitudes across waves

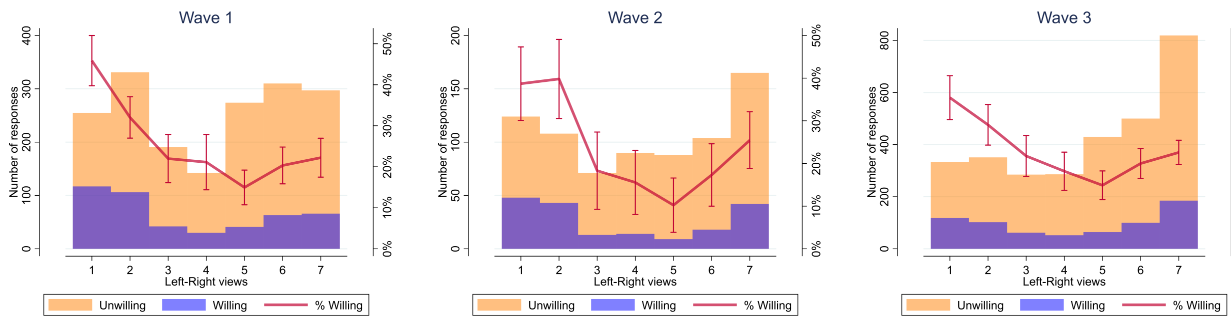

Notes: Agreement to statement in Waves 1-3 by willingness to publish and proportion willing to publish. Responses to Gender statement are reverse-coded such that liberal views are on the left. "Willing" participants are those who respond Yes to being willing to post their opinions on the social media account. Bars represent 95% confidence intervals.

Overall, only 24% of participants were willing to publish their views with their names. Despite a greater number of respondents holding conservative views on the gender and race issues elicited, these individuals were more likely to self-censor than those holding liberal positions, who were most willing to publish (about 40% across the three waves). Moreover, moderates were least willing to publish, while individuals holding more extreme views were relatively more willing to speak publicly, consistent with evidence that political expression on social media is often driven by individuals holding stronger or more extreme views (Bor and Petersen, 2022; Barberá and Rivero, 2015; Bail, 2022; Robertson, Del Rosario and Van Bavel, 2024).

Table 1 presents demographic correlates of willingness to speak up. Across waves, we find that willingness to publish is higher among participants with greater risk tolerance, consistent with the interpretation that publicly expressing one's views is perceived as carrying some degree of social or reputational risk.

[17] Appendix Figure A1 presents distributions by topic, and Appendix Figure A2 by Hi/LoPeer treatment.

Table 1: Correlates of willingness to publish views online with name

	Wave 1				Wave 2				Wave 3			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)	(12)
Stated view (s)		-0.237***		-0.233***		-0.208***		-0.216***		-0.115***		-0.136***
		(0.038)		(0.037)		(0.051)		(0.051)		(0.029)		(0.030)
Extremeness $(2(s-0.5))^2$			0.159***	0.155***			0.167***	0.174***			0.097***	0.114***
			(0.030)	(0.029)			(0.039)	(0.039)			(0.022)	(0.022)
Age	0.004**	0.004**	0.004**	0.004**	0.001	0.002	0.001	0.002*	0.000	0.001	0.000	0.001
	(0.002)	(0.002)	(0.002)	(0.002)	(0.001)	(0.001)	(0.001)	(0.001)	(0.001)	(0.001)	(0.001)	(0.001)
Female	-0.090***	-0.109***	-0.091***	-0.110***	0.017	-0.003	0.014	-0.006	-0.012	-0.017	-0.010	-0.016
	(0.033)	(0.033)	(0.033)	(0.033)	(0.033)	(0.033)	(0.033)	(0.033)	(0.021)	(0.021)	(0.021)	(0.020)
Asian or Pacific Islander	-0.059	-0.057	-0.052	-0.050	-0.120**	-0.108**	-0.096*	-0.083	-0.056	-0.049	-0.051	-0.042
	(0.051)	(0.048)	(0.050)	(0.048)	(0.055)	(0.054)	(0.055)	(0.055)	(0.040)	(0.040)	(0.040)	(0.040)
Black or African American	0.118**	0.120**	0.117**	0.119**	-0.015	0.024	-0.007	0.033	0.024	0.024	0.015	0.014
	(0.051)	(0.050)	(0.051)	(0.050)	(0.049)	(0.049)	(0.048)	(0.049)	(0.040)	(0.040)	(0.041)	(0.041)
Hispanic or Latino	0.021	0.025	0.010	0.015	-0.039	-0.018	-0.017	0.006	0.030	0.036	0.030	0.036
	(0.050)	(0.051)	(0.050)	(0.051)	(0.068)	(0.068)	(0.068)	(0.068)	(0.048)	(0.048)	(0.048)	(0.047)
Other race	-0.034	-0.035	-0.044	-0.045	0.140	0.159	0.153	0.173	-0.006	0.002	-0.021	-0.013
	(0.088)	(0.089)	(0.086)	(0.087)	(0.166)	(0.168)	(0.141)	(0.142)	(0.068)	(0.069)	(0.068)	(0.068)
College degree	-0.025	-0.031	-0.025	-0.032	-0.057	-0.052	-0.046	-0.041	-0.064***	-0.065***	-0.062***	-0.063***
	(0.029)	(0.029)	(0.029)	(0.029)	(0.036)	(0.036)	(0.036)	(0.035)	(0.021)	(0.021)	(0.021)	(0.021)
Employed	0.023	0.025	0.019	0.021	0.078**	0.088**	0.073**	0.083**	0.034	0.037	0.035	0.038*
	(0.030)	(0.029)	(0.030)	(0.029)	(0.037)	(0.036)	(0.037)	(0.036)	(0.023)	(0.023)	(0.023)	(0.022)
Risk attitude	0.082***	0.084***	0.080***	0.082***	0.062***	0.063***	0.061***	0.062***	0.073***	0.073***	0.075***	0.075***
	(0.013)	(0.013)	(0.013)	(0.013)	(0.016)	(0.016)	(0.015)	(0.015)	(0.010)	(0.010)	(0.010)	(0.010)
Active SM users	0.031	0.019	0.026	0.016	0.046	0.036	0.050	0.040	0.070***	0.066***	0.068***	0.062***
	(0.034)	(0.034)	(0.034)	(0.034)	(0.034)	(0.034)	(0.034)	(0.034)	(0.021)	(0.020)	(0.021)	(0.020)
Democrat	0.087***	0.049	0.085***	0.047	0.027	-0.016	0.026	-0.018	0.019	0.002	0.021	0.001
	(0.032)	(0.033)	(0.032)	(0.032)	(0.036)	(0.037)	(0.036)	(0.037)	(0.024)	(0.024)	(0.024)	(0.024)
Republican	-0.054*	0.012	-0.068**	-0.004	-0.044	-0.008	-0.070	-0.034	-0.016	0.012	-0.034	-0.004
	(0.030)	(0.031)	(0.031)	(0.031)	(0.045)	(0.046)	(0.045)	(0.046)	(0.024)	(0.025)	(0.025)	(0.025)
N	1800	1800	1800	1800	750	750	750	750	3004	3004	3004	3004
R-sq	0.087	0.112	0.104	0.129	0.039	0.061	0.061	0.085	0.050	0.057	0.058	0.068

Notes: OLS regressions of willingness to publish with name (0/1) as outcome. *Stated view (s)* is agreement to Gender or Race statement, coded on a 0-1 scale. *Active SM users* indicates the response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

We also correlate willingness to publish with participant's stated view (measured by their agreement to the Gender or Race statement, on a 0-1 scale) and the extremeness of that view (measured by squared deviation from the midpoint, and scaled to 1). Conservative views predict significantly lower willingness to publish, even after controlling for covariates and party identification, suggesting this pattern is not solely driven by party-level differences in public expression. Furthermore, conditional on controls, those holding extreme views also have a higher willingness to publish. Using the estimates in columns 4, 8 and 12 respectively, willingness to speak up is bimodal at the extremes but is highest at the left-most position, consistent with the patterns in Figure 5.[18]

To investigate whether social pressure and perceived issue importance can explain these patterns, we add further controls for respondents' self-censoring, using their response to "The political climate these days prevents me from saying things I believe because others might find them offensive" (1 Strongly disagree - 7 Strongly agree), and topic importance, which measures how important respondents consider the issue discussed in the statement. The results are presented in Appendix Table A7. Across all waves, controlling for self-censoring substantially attenuates the negative relationship between conservative views and willingness to publish, while leaving the extremeness coefficient largely unchanged. This finding suggests that conservatives are less willing to publish their views partly because they perceive greater social and reputational costs from doing so. By contrast, controlling for topic importance mainly reduces the coefficient on extremeness, indicating that individuals with more extreme views are more willing to publish partly because they care more strongly about the issues themselves. When controlling for both self-censoring and topic importance, both coefficients are reduced, and in Wave 3 the coefficient on stated view is no longer statistically significant. Furthermore, we plot these two variables against ideological stance in Figure A3, which reveals a mostly linear relationship between ideology and participants' self-reported self-censorship due to the political climate, and a U-shaped relationship between ideology and topic importance. Together, these results suggest that the two patterns have distinct drivers: conservatives' reluctance to publish reflects perceived social costs, while the greater willingness of those with extreme views reflects stronger engagement with the issues.

In Appendix A.1, we discuss an extension of our model in which the intrinsic payoff from speaking up increases with the extremeness of private opinions.

[18]The estimates imply that willingness to speak up is maximized at $s = 0$ and minimized at $s = 0.69, 0.66, 0.65$ (on a 0-1 scale). Our results are robust to controlling for frequency of posting, see Table A5. Similar patterns hold using our incentivized measure of willingness to pay to publish (Appendix Table A6).

Beliefs about others' views

We elicited participants' incentivized beliefs about others' views, distinguishing between beliefs about (i) all participants in the study and (ii) participants willing to publish their views (without payment). We define the *perceived gap* as $b_{all} - b_{pub}$, where b_{all} represents a participant's belief about the average (or majority, in Wave 3) view of all participants in the wave, and b_{pub} reflects their belief about the views of those willing to publish. This gap captures the extent to which participants expect publicly expressed views to differ from the broader distribution of beliefs. The results are presented in Table 2 by wave and topic.

Participants expected publicly expressed views to be more progressive than the overall distribution of opinions. Specifically, respondents believed that the views of participants willing to publish would lie to the left of the average view among all participants. The perceived gap between b_{all} and b_{pub} is largest and statistically significant for the race topic, while the gender topic shows the same directional pattern, but is only consistently significant in Wave 3. Within the framework of our model, these beliefs are consistent with a progressive expressive norm: views closer to the progressive end of the scale being perceived as less socially costly to express publicly.

4.2 Conformity: Treatment effects on reported attitudes

We first analyze whether participants changed their stated views (s) in response to the Prime and Awareness treatments, as tested in Wave 1 of Experiment 1. We define s_{iq} as respondent i's reported stance (scaled from 0 to 1) for question q and estimate two pre-registered specifications:

$$s_{iq} = \theta_0 + \theta_1 Awareness_i + \theta_2 Prime_i + \delta_q + \varepsilon_{iq}$$
$$s_{iq} = \theta_0 + \theta_1 Awareness_i + \theta_2 Prime_i + \theta_3 Awareness_i \times Prime_i + \delta_q + \varepsilon_{iq}$$

where $Prime_i$ equals 1 if subject i sees the online backlash prime before responding (T1 and T4), $Awareness_i$ equals 1 if the response is elicited when subjects are aware of potential publication (T4 and T5), and ε_{iq} is an individual-question specific error term. We include topic fixed effects δ_q. In some specifications we include a vector of pre-registered controls $\mathbf{X_i}$ including age, gender, race, education, employment, risk attitude, political leaning, and social media use, which may increase the precision of our estimates (but are by construction uncorrelated with our randomized treatment). Robust standard errors are clustered at the individual level (since we have two observations per subject).

Table 3 presents the results. We first examine the average effect of publication awareness in columns 1-2. The coefficients are small and statistically indistinguishable from

Table 2: Summary statistics on beliefs about others' views.

	b_{all}	b_{pub}	$b_{all} - b_{pub}$
Wave 1, gender topic			
All	3.882	3.744	0.138
Cons	4.272	4.111	0.160
Mod	3.817	3.754	0.063
Prog	3.540	3.366	0.174
Wave 1, race topic			
All	4.499	3.545	0.953***
Cons	4.887	4.060	0.827***
Mod	4.495	3.468	1.027***
Prog	4.114	3.136	0.977***
Wave 2, gender topic			
All	4.027	3.901	0.125*
Cons	4.379	4.190	0.190
Mod	4.008	3.928	0.080
Prog	3.638	3.539	0.099
Wave 3, gender topic			
All	4.868	4.050	0.819***
Cons	5.296	4.337	0.958***
Mod	4.495	3.768	0.727***
Prog	4.101	3.588	0.513***
Wave 3, race topic			
All	4.517	3.590	0.926***
Cons	5.103	3.853	1.250***
Mod	4.583	3.816	0.767***
Prog	3.766	3.005	0.761***

Notes: Belief about views of all participants in the study, all participants willing to publish, and the difference between these two values. In Experiment 1 (Waves 1 and 2) we elicit beliefs about the *average* view of other participants, in Experiment 2 (Wave 3) we elicit beliefs about the *majority* view of other participants. The gender question is reverse-coded: 1 = progressive, 7 = conservative. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

zero, indicating no clear overall shift in stated views. In columns 3-4, we allow the effect of awareness to vary with the prime. In the absence of the prime, publication awareness shifts responses by about 0.051 on the 0-1 scale in the progressive direction (column 4).[19]

Table 3: Conformity in stated views (Wave 1)

	All		All		Q1		Q2	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Awareness	-0.001	-0.010	-0.041	-0.051**	-0.019	-0.031	-0.061*	-0.071**
	(0.022)	(0.016)	(0.030)	(0.022)	(0.035)	(0.026)	(0.035)	(0.031)
Prime	0.000	-0.002	-0.025	-0.028	-0.002	-0.003	-0.047*	-0.052**
	(0.020)	(0.015)	(0.025)	(0.018)	(0.029)	(0.023)	(0.028)	(0.023)
Prime x Awareness			0.082*	0.084***	0.082	0.084**	0.081	0.084*
			(0.043)	(0.032)	(0.051)	(0.040)	(0.050)	(0.043)
Mean in control	0.527	0.527	0.527	0.527	0.514	0.514	0.539	0.539
N	1800	1800	1800	1800	900	900	900	900
R-sq	0.001	0.357	0.004	0.360	0.008	0.400	0.005	0.333
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of stated view s, defined as agreement to topic statement (reverse-coded for Gender statement) and scaled to 0-1. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Columns 5-8 split the analysis by question in the survey order. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Turning to the Prime intervention, the average effect is small and imprecisely estimated (row 2 of Table 3). While the point estimates are generally negative, suggesting a shift toward more progressive responses, they are not statistically significant. The effects are more pronounced, however, for the second question (when participants are aware of the possibility of publication) and among conservatives (Table A8).

Finally, the interaction between the prime and publication awareness runs counter to our pre-registered hypothesis. The interaction term is positive and statistically significant (row 3 of Table 3), implying that the prime attenuates, rather than amplifies, the shift toward more progressive responses induced by publication awareness. One possible interpretation is that, when participants anticipate publication, the prime increases the salience of competing considerations—such as free speech concerns or resistance to

[19]The effect is driven by responses to the second question, which we present as a complementary exploratory analysis (Table 3, columns 5-8). This pattern is consistent with the design, since respondents in the awareness condition are explicitly informed of potential publication after answering the first question but before answering the second. We do not observe a significant effect for the first question (columns 5-6). Only six subjects used the "Back" button to revise their initial response to the first question. In Table A8, we explore heterogeneity by ideological stance, using responses to the first question to classify participants and responses to the second question as the outcome. The effects are concentrated among participants with moderate views and, to a lesser extent, those with conservative views.

perceived pressure—though the experiment was not designed to isolate this mechanism. Notably, the combined effect of the prime and publication awareness (the sum of rows 1 and 3) is small and statistically indistinguishable from zero.

4.3 Silence: Treatment effects on public expression

Baseline effects

We next study individuals' willingness to publish their opinion with their name (v). We regress the outcome variable on the relevant treatment dummies as follows:

$$v_{iq} = \theta_0 + \theta_1 Prime_i + \theta_2 T_i + \delta_q + \varepsilon_{iq}$$

where v_{iq} is a binary variable equal to 1 if subject i is willing to publish their opinion on question q with their name (without additional lottery tickets as incentive). Depending on the wave, the alternative treatment T_i indicates whether the participants were in the Awareness (Wave 1) or the HiPeer (Wave 3) treatment groups. The parameters θ_1 and θ_2 capture the average treatment effect of our interventions. We also include topic fixed effects δ_q. In some specifications we include additional demographic controls.

Additionally, we split the sample according to participants' privately stated views to examine whether treatment effects vary across ideological positions. Responses are coded so that lower values correspond to more progressive views. We classify respondents as Progressive if their stated view is 1 or 2 (on the 1–7 scale), Conservative if it is 6 or 7, and Moderate otherwise (3, 4, or 5).

The results (Table 4) show that, on average, neither the online backlash Prime nor the Publication Awareness treatment significantly affect willingness to publish in the first set of experimental waves (Waves 1-2). In Wave 3, conducted two years later, the Prime reduces willingness to publish. However, we find no significant effect overall when pooling across Waves 1-3. This effect is also not statistically significant for participants who were not exposed to the HiPeer treatment (see section below on interacted effects), consistent with the findings from the previous two waves.

The HiPeer treatment—in which respondents learn that a high share of previous participants published their opinions—has a statistically significant positive effect, increasing willingness to speak up by about 5 percentage points relative to those in the LoPeer group (and shown in Appendix Figure A2 by reported attitude).[20]

[20] This effect is smaller than in similar interventions, such as Bursztyn et al. (2023), where social cover increased public expression by 11 percentage points among progressives. Additionally, we pre-registered the WTP to publish as an alternative outcome variable. The results, displayed in Appendix Table A9, are

Table 4: Willingness to publish and experimental interventions

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
				Wave 1				
Prime	0.044	0.051*	0.087	0.094*	0.042	0.054	-0.003	0.004
	(0.029)	(0.028)	(0.053)	(0.054)	(0.040)	(0.039)	(0.043)	(0.040)
Awareness	0.020	0.029	0.040	0.047	-0.001	0.002	0.025	0.052
	(0.030)	(0.029)	(0.054)	(0.054)	(0.041)	(0.042)	(0.043)	(0.040)
Mean in control	0.213	0.213	0.303	0.303	0.151	0.151	0.204	0.204
N	1800	1800	586	586	607	607	607	607
R-sq	0.002	0.106	0.015	0.079	0.003	0.108	0.014	0.178
				Wave 2				
Prime	-0.017	-0.015	-0.053	-0.056	0.001	0.006	-0.005	-0.018
	(0.032)	(0.031)	(0.064)	(0.066)	(0.045)	(0.044)	(0.051)	(0.050)
Mean in control	0.258	0.258	0.420	0.420	0.144	0.144	0.225	0.225
N	750	750	232	232	249	249	269	269
R-sq	0.000	0.059	0.003	0.059	0.000	0.123	0.000	0.085
				Wave 3				
Prime	-0.058***	-0.048**	-0.024	-0.028	-0.064**	-0.055*	-0.075**	-0.066**
	(0.022)	(0.022)	(0.046)	(0.045)	(0.030)	(0.030)	(0.031)	(0.030)
HiPeer	0.054***	0.044**	0.064	0.044	0.001	0.000	0.091***	0.069**
	(0.020)	(0.020)	(0.043)	(0.041)	(0.028)	(0.028)	(0.028)	(0.027)
Mean in control	0.230	0.230	0.333	0.333	0.205	0.205	0.202	0.202
N	3004	3004	684	684	1001	1001	1319	1319
R-sq	0.009	0.062	0.021	0.102	0.006	0.057	0.019	0.064
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to publish with name (0/1). Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. We classify respondents as Progressive if their stated view is 1 or 2 (on the 1–7 scale left-right scale), Conservative if it is 6 or 7, and Moderate otherwise (3, 4, or 5). Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Notably, the effects of the Prime and HiPeer treatments are largest among participants with conservative views, as shown in columns 7-8. This heterogeneity is consistent with the model's prediction that treatment effects should be larger for participants whose views are farther from the socially accepted expressive norm. The heterogeneity may reflect changes in perceived costs of political expression and resulting strategic responses, although our design does not allow us to isolate the underlying mechanism (as discussed below).

Interacted effects

Next, we explore pre-registered empirical extensions which interact our Prime treatment with the alternative treatment. For Wave 1, we interact Prime with Awareness to test whether the Prime's effect is attenuated when opinions may adjust due to potential publication. For Wave 3, we interact the Prime treatment with HiPeer to examine whether the effect of high peer participation found above varies with social pressure. The models take the following form:

$$v_{iq} = \theta_0 + \theta_1 Prime_i + \theta_2 T_i + \theta_3 Prime_i \times T_i + \delta_q + \varepsilon_{iq} \tag{6}$$

where, as before, $T_i \in \{Awareness_i, HiPeer_i\}$.

The results for Wave 1 are shown in Table 5. We hypothesized that the online backlash Prime would have a stronger negative effect when individuals are unaware of potential publication, as those in the Publication Awareness treatment could adjust s to conform. However, we do not find a statistically significant interaction between the Prime and Awareness treatments.

Table 6 presents the results for Wave 3. The HiPeer treatment generally increases willingness to publish ($\theta_2 > 0$), but this effect is concentrated among participants who were not exposed to the online backlash prime. In other words, the effect of HiPeer is attenuated in the presence of the prime ($\theta_3 < 0$), though the interaction itself is not statistically significant. The combined effect ($\theta_2 + \theta_3$) is not statistically different from zero. One interpretation consistent with this pattern is that higher perceived costs of political expression increase free-riding incentives among individuals holding dissenting or conservative views. While observing others speak up may create strategic complementarities that encourage expression, these incentives appear to be offset when the perceived risk of backlash increases,

very similar to the main results presented above, though less statistically precise. We perform two additional wave-specific robustness checks. In Wave 1, to address potential endogeneity in reported attitudes (s), we exclude treatments where attitudes are measured post-treatment. The results are very similar and shown in Appendix Table A10. Controlling for posting frequency (Wave 3) also does not qualitatively change our results, see Table A11.

Table 5: Willingness to publish and experimental interventions (Wave 1)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	0.055	0.070**	0.093	0.114	0.051	0.065	-0.002	0.016
	(0.037)	(0.035)	(0.068)	(0.070)	(0.051)	(0.046)	(0.055)	(0.052)
Awareness	0.037	0.059	0.048	0.077	0.013	0.020	0.028	0.072
	(0.046)	(0.045)	(0.082)	(0.083)	(0.059)	(0.058)	(0.072)	(0.068)
Prime x Awareness	-0.028	-0.050	-0.013	-0.049	-0.024	-0.030	-0.004	-0.032
	(0.061)	(0.059)	(0.109)	(0.110)	(0.081)	(0.081)	(0.090)	(0.083)
Mean in control	0.213	0.213	0.303	0.303	0.151	0.151	0.204	0.204
N	1800	1800	586	586	607	607	607	607
R-sq	0.002	0.106	0.015	0.079	0.003	0.108	0.014	0.178
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to publish with name (0/1). Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. We classify respondents as Progressive if their stated view is 1 or 2 (on the 1–7 scale left-right scale), Conservative if it is 6 or 7, and Moderate otherwise (3, 4, or 5). Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

leading individuals to rely on others to voice dissent.[21]

While our experimental design does not pinpoint the exact mechanism behind the HiPeer treatment's effect on willingness to publish, we explore potential explanations in Section 5.1 below.

4.4 Heterogeneous treatment effects

We explore whether the effect of the Prime treatment was heterogeneous along various dimensions with specifications of the following form:

$$v_{iq} = \theta_0 + \theta_1 Prime_i + \theta_2 Var_i + \theta_3 Prime_i \times Var_i + \delta_q + \varepsilon_{iq}$$

where Var_i indicates the dimension of interest.[22] In Wave 1 we include the Awareness treatment dummy while in Wave 3 we include the HiPeer treatment dummy. In Figure 6 we report the coefficients θ_3 for these interaction models.

The variables interacted with our treatments include risk attitudes, active social media use, topic importance, stated view, an index of psychological reactance (Wave 3 only), perceived popularity of first and last names (Wave 3 only), and other demographic charac-

[21]We replicate these analyses for willingness to pay in Appendix Tables A13 and A14. Controlling for posting frequency (Wave 3) also does not qualitatively change our results, see Table A11.

[22]Heterogeneity analyses for the Publication Awareness and HiPeer treatments are shown in Appendix Figures A4 and A5 respectively.

Table 6: Willingness to publish and experimental interventions (Wave 3)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	-0.044	-0.025	-0.062	-0.044	-0.042	-0.037	-0.043	-0.024
	(0.029)	(0.029)	(0.063)	(0.062)	(0.043)	(0.044)	(0.038)	(0.038)
HiPeer	0.074**	0.075**	0.011	0.022	0.030	0.023	0.138***	0.130**
	(0.037)	(0.036)	(0.075)	(0.075)	(0.052)	(0.052)	(0.053)	(0.052)
Prime x HiPeer	-0.029	-0.046	0.077	0.032	-0.044	-0.036	-0.068	-0.090
	(0.044)	(0.043)	(0.091)	(0.091)	(0.061)	(0.062)	(0.062)	(0.061)
Mean in control	0.230	0.230	0.333	0.333	0.205	0.205	0.202	0.202
N	3004	3004	684	684	1001	1001	1319	1319
R-sq	0.009	0.063	0.022	0.102	0.007	0.058	0.021	0.066
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to publish with name (0/1). Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. We classify respondents as Progressive if their stated view is 1 or 2 (on the 1–7 scale left-right scale), Conservative if it is 6 or 7, and Moderate otherwise (3, 4, or 5). Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

teristics. With the exception of the dummy variables *Female*, *College*, *Employed*, and *Active social media use*, all variables are standardized to facilitate interpretation. All of these dimensions of heterogeneity are of interest but we were generally agnostic as to the potential direction of the effects (as outlined in our pre-registration).

We find statistically significant heterogeneity along *Reactance* and *Age* in Wave 3. Participants with higher reactance exhibit a backlash to the Prime treatment, increasing their willingness to speak up, whereas those with lower reactance exhibit greater self-censorship in response to the treatment. A similar pattern emerges for age: older participants become less willing to publish their views following the prime, while younger participants become more willing to speak up. That certain types of less/more inhibited individuals respond differentially to the Prime treatment is perhaps not surprising, and it suggests that public views may be disproportionately shaped by vocal minorities who feel less constrained by social repercussions.

5 Empirical Extensions

5.1 Mechanisms in HiPeer treatment

The positive effect of peer participation, captured by the HiPeer treatment, may operate through several channels, including changes in perceived social pressure, beliefs about oth-

Figure 6: Heterogeneity of Prime treatment

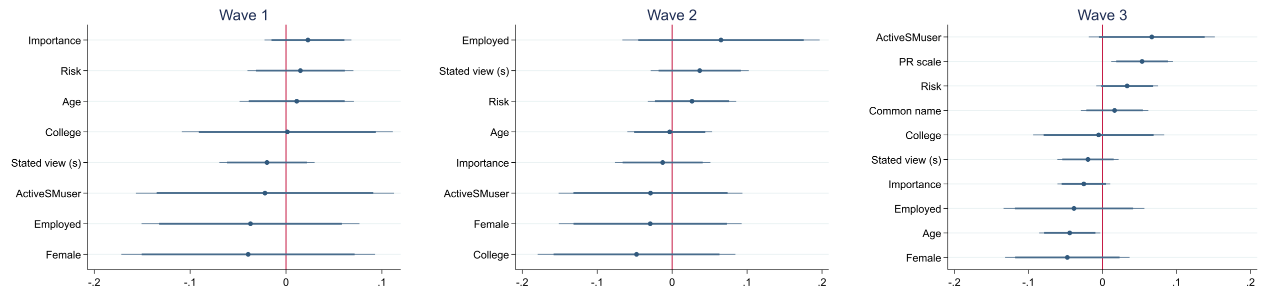

Notes: The figure shows the coefficient θ_3 from the following model:

$$v_{iq} = \theta_0 + \theta_1 Prime_i + \theta_2 Var_i + \theta_3 Prime_i \times Var_i + \delta_q + \varepsilon_{iq}$$

where Var_i indicates the dimension of interest in exploring heterogeneous effects. In Wave 1 we include the Awareness treatment dummy while in Wave 3 we include the HiPeer treatment dummy. *Stated view (s):* agreement to topic statement (reverse-coded for Gender statement) and scaled to 0-1. *Risk:* response to "Please tell us, in general, how willing or unwilling you are to take risks." (0 Completely unwilling to take risks - 10 Very willing to take risks). *Active SM user:* response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour. *Importance:* response to "How important is the issue discussed in the statement to you?" (1 Not important at all - 5 Extremely important). *Common name:* sum of responses to "The most common first name in the US is James. How common do you think your first name is in the US?" and "The most common last name in the US is Smith. How common do you think your last name is in the US?" (0 Extremely uncommon - 7 Extremely common, Wave 3 only). *PR Scale:* responses to the Hong psychological reactance scale (Hong and Faedda, 1996) (Wave 3 only). Except for the dummy variables Female, College, Employed, Active social media use, all other variables are standardised. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors are clustered at the individual level.

ers' behavior, or other responses to the information provided. In this section, we examine these possibilities descriptively. The analyses below were not part of our pre-registration and should therefore be interpreted as exploratory.

First, the effect of the HiPeer treatment on willingness to speak up is concentrated among participants who did not receive the Prime treatment. One possible interpretation is that participants who are less aligned with perceived publicly expressed views are more responsive to information about others' willingness to publish. As suggestive evidence, Appendix Table A16 reports the effect of the HiPeer treatment separately for progressives, moderates, and conservatives, splitting the sample by Prime status. We find that the HiPeer effect is largest for respondents with more conservative views who were not primed to consider online backlash. While this pattern is consistent with changes in the perceived social costs of expression, it may also reflect other mechanisms, and the design does not allow us to distinguish between them.

To examine whether the HiPeer treatment affects beliefs about expressed views, we use our measures of b_{all}, the perceived opinion of all participants, and b_{pub}, the perceived opinion of participants willing to publish. As described above, we transform these measures to a 0–1 left-right attitude scale. We regress these outcomes on the HiPeer and Prime treatment indicators and their interaction. The results are shown in Appendix Table A17.

As shown in columns 1–4, the HiPeer treatment does not significantly affect beliefs about the opinions of all other participants. However, being informed that many others were willing to publish shifts beliefs about the views of participants willing to publish to the right (columns 5-8). As a result, the perceived gap between all participants' views and publicly expressed views, measured by $b_{all} - b_{pub}$, is significantly smaller in the HiPeer treatment (columns 9-12). These results show that the HiPeer treatment changed participants' beliefs about which views were likely to be publicly expressed, but they do not uniquely identify the channel through which the treatment affected willingness to publish.

5.2 Open text analysis

At the end of Wave 3, we asked participants in an open text box about their reasons for (agreeing to) posting, or not, their opinions on Twitter. We classify these responses using a large language model (LLM), as outlined in Appendix Section A.3. We then analyzed the classified responses, and the results are presented below. Although not pre-registered, this analysis provides additional insight into how social pressure influences public expression in our survey.

Descriptive evidence of reasons for (not) posting

Appendix Figure A6 shows the categorization of reasons given by participants for posting/not posting opinions across the ideological spectrum of the stated opinion. We observe some key empirical patterns which are consistent with our results above. First, among those not posting, fear of backlash is more prevalent on the right—consistent with the view that left-wing norms tend to constrain expression. Second, freedom of speech is more frequently cited by those on the right as a reason for posting. Third, viewing the issue as important is a common reason among individuals at the ends of the ideological distribution for posting their views (somewhat more on left), and conversely, indifference to the issues appears as a relatively more common reason among moderates for not posting, consistent with our evidence above of topic importance explaining the relationship between speaking up and ideological extremity. This final observation highlights one important driver of public expression: the degree to which individuals care about these particular issues.

Regression analysis

Appendix Table A18 shows OLS regressions of the likelihood of stating a particular reason for posting/not posting on our treatment indicators and the ideological spectrum as coded by participants' opinions. Indifference to social pressure is more frequently cited as a reason for posting among those in the HiPeer group. For those unwilling to post, fear of backlash is more likely to be stated as participants' views move to the right, consistent with the existence of a liberal expressive norm. Interest in and importance of the race and gender issues are less frequent for moderates, and somewhat less salient to those in the HiPeer group.

Appendix Table A19 shows results for the other reasons. Freedom of speech is more likely mentioned for those in the centre and more right-wing participants, though the relationship becomes statistically insignificant with controls. Those on the right are less likely to cite privacy concerns but more likely to care about financial incentives. They are also less likely to cite disengagement from social media as a reason for not posting.

Keyword analysis

Appendix Figure A7 shows the proportion of respondents mentioning certain keywords related to our specific research hypotheses, and how these vary across the left-right ideology scale. Keywords related to social backlash are more common among right-leaning participants, keywords related to race or gender tend to be least common among moderate participants (and most common among left-leaning participants), and keywords related to other people increase with left-right ideology.

Treatment effects mediators

Appendix Tables A20 and A21 regress the number of statements published (0, 1 or 2) on our main treatment indicators (HiPeer or Prime) in Wave 3, and controlling in turn for each possible reason to post/not to post. This exercise includes each reason as a "bad control" in a baseline regression framework to understand whether specific reasons may mediate the treatment effect. The effect of the HiPeer treatment becomes statistically insignificant when controlling for "indifference to social pressure", and when controlling for "privacy concerns". The effect of the Prime treatment is attenuated when controlling for "indifference to social pressure" and "freedom of speech".

Altogether, our open text analyses provide evidence consistent with our other findings. Social backlash is most salient for those on the right; however, the HiPeer treatment reduces perceptions of social pressure, while the Prime has the opposite effect. Importantly, willingness to speak up is shaped not only by fear of backlash but also by how much individuals value the issues.

6 Conclusion

Our experiment provides insights into online public expression in an environment of growing social and political polarization. We find that a substantial share of individuals are hesitant to share their views on polarizing topics, and many of them cite fear of social backlash as an important reason for this hesitancy. Although conservative attitudes were more common in our sample, participants with left-leaning views were more willing to speak publicly. Participants also perceived publicly expressed views to be more progressive than the overall distribution of opinions. Taken together, these patterns suggest that progressive positions were perceived as less socially costly to voice, consistent with the presence of a left-leaning expressive norm in this setting. Furthermore, we show that individuals holding moderate views were the least willing to publish their opinion in our experiment, and they report viewing these issues as less important. We further show that priming participants to consider the risk of negative social backlash had only modest effects on willingness to speak up.

Free speech is a key component of liberal democracies, and many participants in our experiment reported fear of repercussions as a reason not to voice their views publicly. At the same time, for better or worse, strong norms which induce self-censorship can unravel quickly (Morales, 2020; Bursztyn, Egorov and Fiorin, 2020; Álvarez-Benjumea, 2023; Albornoz, Bradley and Sonderegger, 2022; Apffelstaedt, Freundt and Oslislo, 2022). We provide evidence that informing people about others' willingness to speak up can reduce self-censorship, suggesting that simple interventions may increase people's willingness to share their views and shift perceptions around public speech norms. Interestingly, this result was present despite participants not being informed about *what* others were sharing, suggesting that a norm of free speech more generally may play a role in shaping online behavior, and not just a norm of what constitutes acceptable speech.

While our studies were conducted at a time when liberal norms were especially salient in both social and mainstream media, the pressure to conform to group norms is a general phenomenon that is likely to persist even as the specific norms themselves shift. Understanding how social pressure and group dynamics shape online public expression, and how this in turn influences discourse on sensitive social and political issues, remains an important topic of study in a persistently polarized media environment.

Declaration

Declaration of generative AI and AI-assisted technologies in the writing process. During the preparation of this work the authors used ChatGPT in order to improve the readability and language of some sentences in the manuscript. After using this tool/service, the authors reviewed and edited the content as needed and take full responsibility for the content of the published article.

References

Acemoglu, Daron, Georgy Egorov, and Konstantin Sonin. 2026. "Multidimensional signaling and the rise of cultural politics." *NBER Working Paper 34909*.

Akerlof, George A, and Rachel E Kranton. 2000. "Economics and identity." *The Quarterly Journal of Economics*, 115(3): 715–753.

Albornoz, Facundo, Jake Bradley, and Silvia Sonderegger. 2022. "Updating the social norm: The case of hate crime after the Brexit referendum." Working Paper.

Álvarez-Benjumea, Amalia. 2023. "Uncovering hidden opinions: Social norms and the expression of xenophobic attitudes." *European Sociological Review*, 39(3): 449–463.

Andreoni, James, and B Douglas Bernheim. 2009. "Social image and the 50–50 norm: A theoretical and experimental analysis of audience effects." *Econometrica*, 77(5): 1607–1636.

Andreoni, James, and Ragan Petrie. 2004. "Public goods experiments without confidentiality: A glimpse into fund-raising." *Journal of Public Economics*, 88(7-8): 1605–1623.

Apffelstaedt, Arno, Jana Freundt, and Christoph Oslislo. 2022. "Social norms and elections: How elected rules can make behavior (in) appropriate." *Journal of Economic Behavior & Organization*, 196: 148–177.

Bail, Chris. 2022. *Breaking the social media prism: How to make our platforms less polarizing.* Princeton University Press.

Barberá, Pablo, and Gonzalo Rivero. 2015. "Understanding the political representativeness of Twitter users." *Social Science Computer Review*, 33(6): 712–729.

Bašić, Zvonimir, and Eugenio Verrina. 2024. "Personal norms—and not only social norms—shape economic behavior." *Journal of Public Economics*, 239: 105255.

Bénabou, Roland, and Jean Tirole. 2006. "Incentives and prosocial behavior." *American Economic Review*, 96(5): 1652–1678.

Benson, Jonathan. 2024. "Democracy and the epistemic problems of political polarization." *American Political Science Review*, 118(4): 1719–1732.

Bernheim, B Douglas. 1994. "A theory of conformity." *Journal of Political Economy*, 102(5): 841–877.

Bond, Robert M, Christopher J Fariss, Jason J Jones, Adam DI Kramer, Cameron Marlow, Jaime E Settle, and James H Fowler. 2012. "A 61-million-person experiment in social influence and political mobilization." *Nature*, 489(7415): 295–298.

Bor, Alexander, and Michael Bang Petersen. 2022. "The psychology of online political hostility: A comprehensive, cross-national test of the mismatch hypothesis." *American Political Science Review*, 116(1): 1–18.

Boxell, Levi, Matthew Gentzkow, and Jesse M Shapiro. 2024. "Cross-country trends in affective polarization." *Review of Economics and Statistics*, 106(2): 557–565.

Brady, William J., Killian McLoughlin, Tuan N. Doan, and Molly J. Crockett. 2021. "How social learning amplifies moral outrage expression in online social networks." *Science Advances*, 7(33): eabe5641.

Braghieri, Luca. 2024. "Political correctness, social image, and information transmission." *American Economic Review*, 114(12): 3877–3904.

Braghieri, Luca, Leonardo Bursztyn, and Jan Fasnacht. 2026. "Threshold disclosure in collective decisions." Working Paper.

Bursztyn, Leonardo, and Robert Jensen. 2017. "Social image and economic behavior in the field: Identifying, understanding, and shaping social pressure." *Annual Review of Economics*, 9: 131–153.

Bursztyn, Leonardo, Georgy Egorov, and Stefano Fiorin. 2020. "From extreme to mainstream: The erosion of social norms." *American Economic Review*, 110(11): 3522–48.

Bursztyn, Leonardo, Georgy Egorov, Ingar Haaland, Aakaash Rao, and Christopher Roth. 2023. "Justifying dissent." *The Quarterly Journal of Economics*, 138(3): 1403–1451.

Cantoni, Davide, David Y Yang, Noam Yuchtman, and Y Jane Zhang. 2019. "Protests as strategic games: Experimental evidence from Hong Kong's antiauthoritarian movement." *The Quarterly Journal of Economics*, 134(2): 1021–1077.

Carlson, Taylor N, and Jaime E Settle. 2016. "Political chameleons: An exploration of conformity in political discussions." *Political Behavior*, 38(4): 817–859.

Conzo ⓡ, Pierluigi, Laura K Taylor ⓡ, Juan S Morales ⓡ, Margaret Samahita ⓡ, and Andrea Gallice. 2023. "Can ♥s change minds? Social media endorsements and policy preferences." *Social Media + Society*, 9(2): 20563051231177899.

Crockett, Molly J. 2017. "Moral outrage in the digital age." *Nature Human Behaviour*, 1(11): 769–771.

DellaVigna, Stefano, John A List, Ulrike Malmendier, and Gautam Rao. 2016. "Voting to tell others." *The Review of Economic Studies*, 84(1): 143–181.

De Quidt, Jonathan, Johannes Haushofer, and Christopher Roth. 2018. "Measuring and bounding experimenter demand." *American Economic Review*, 108(11): 3266–3302.

Dias, Nicholas C, James N Druckman, and Matthew S Levendusky. 2025. "Unraveling a "Cancel Culture" dynamic: When, why, and which Americans sanction offensive speech." *The Journal of Politics*, 87(2): 588–600.

Dimant, Eugen. 2024. "Hate trumps love: The impact of political polarization on social preferences." *Management Science*, 70(1): 1–31.

Dinas, Elias, Sergi Martínez, and Vicente Valentim. 2024. "Social norm change, political symbols, and expression of stigmatized preferences." *The Journal of Politics*, 86(2): 488–506.

Draca, Mirko, and Carlo Schwarz. 2024. "How polarised are citizens? Measuring ideology from the ground up." *The Economic Journal*, 134(661): 1950–1984.

Forestal, Jennifer. 2024. "Social media, social control, and the politics of public shaming." *American Political Science Review*, 118(4): 1704–1718.

Friedrichsen, Jana, and Dirk Engelmann. 2018. "Who cares about social image?" *European Economic Review*, 110: 61–77.

Gibson, James L, and Joseph L Sutherland. 2023. "Keeping your mouth shut: Spiraling self-censorship in the United States." *Political Science Quarterly*, 138(3): 361–376.

Golman, Russell. 2023. "Acceptable discourse: Social norms of beliefs and opinions." *European Economic Review*, 160: 104588.

González, Felipe. 2020. "Collective action in networks: Evidence from the Chilean student movement." *Journal of Public Economics*, 188: 104220.

Hager, Anselm, Lukas Hensel, Johannes Hermle, and Christopher Roth. 2023. "Political activists as free riders: Evidence from a natural field experiment." *The Economic Journal*, 133(653): 2068–2084.

Hauk, Esther, and Javier Ortega. 2025. "Political correctness and elite prestige." *Journal of Theoretical Politics*, 37(2): 122–155.

Holm, Hakan J, and Margaret Samahita. 2018. "Curating social image: Experimental evidence on the value of actions and selfies." *Journal of Economic Behavior & Organization*, 148: 83–104.

Hong, Sung-Mook, and Salvatora Faedda. 1996. "Refinement of the Hong psychological reactance scale." *Educational and Psychological Measurement*, 56(1): 173–182.

Ho, Yuen, and Yihong Huang. 2024. "Breaking the spiral of silence." Working Paper.

Isler, Ozan, and Simon Gächter. 2022. "Conforming with peers in honesty and cooperation." *Journal of Economic Behavior & Organization*, 195: 75–86.

Kimbrough, Erik O, and Alexander Vostroknutov. 2016. "Norms make preferences social." *Journal of the European Economic Association*, 14(3): 608–638.

Krupka, Erin L, and Roberto A Weber. 2013. "Identifying social norms using coordination games: Why does dictator game sharing vary?" *Journal of the European Economic Association*, 11(3): 495–524.

Kuran, Timur. 1997. *Private truths, public lies.* Harvard University Press.

Le Yaouanq, Yves, Peter Schwardmann, and Joël J van der Weele. 2025. "Rationalizations and political polarization." CESifo Working Paper No. 11897.

Manacorda, Marco, and Andrea Tesei. 2020. "Liberation technology: Mobile phones and political mobilization in Africa." *Econometrica*, 88(2): 533–567.

Michaeli, Moti, and Daniel Spiro. 2017. "From peer pressure to biased norms." *American Economic Journal: Microeconomics*, 9(1): 152–216.

Montgomery, Jacob M, Brendan Nyhan, and Michelle Torres. 2018. "How conditioning on posttreatment variables can ruin your experiment and what to do about it." *American Journal of Political Science*, 62(3): 760–775.

Morales, Juan S. 2020. "Perceived popularity and online political dissent: Evidence from Twitter in Venezuela." *The International Journal of Press/Politics*, 25(1): 5–27.

Morris, Stephen. 2001. "Political correctness." *Journal of Political Economy*, 109(2): 231–265.

Mummolo, Jonathan, and Erik Peterson. 2019. "Demand effects in survey experiments: An empirical assessment." *American Political Science Review*, 113(2): 517–529.

Noelle-Neumann, Elisabeth. 1974. "The spiral of silence: A theory of public opinion." *Journal of Communication*, 24(2): 43–51.

Norris, Pippa. 2023. "Cancel culture: Myth or reality?" *Political Studies*, 71(1): 145–174.

Norris, Pippa. 2025. "Cancel culture: Heterodox self-censorship or the curious case of the dog which didn't bark." *International Political Science Review*, 46(3): 422–441.

Perez-Truglia, Ricardo, and Guillermo Cruces. 2017. "Partisan interactions: Evidence from a field experiment in the united states." *Journal of Political Economy*, 125(4): 1208–1243.

Rege, Mari, and Kjetil Telle. 2004. "The impact of social approval and framing on cooperation in public good situations." *Journal of Public Economics*, 88(7-8): 1625–1644.

Robertson, Claire E, Kareena S Del Rosario, and Jay J Van Bavel. 2024. "Inside the funhouse mirror factory: How social media distorts perceptions of norms." *Current Opinion in Psychology*, 60: 101918.

Schulz, William Small. 2025. "What would you say? Estimating causal effects of social context on political expression." *Political Analysis*, 33(2): 122–139.

Valentim, Vicente. 2024. "Political stigma and preference falsification: theory and observational evidence." *The Journal of Politics*, 86(4): 1382–1402.

Voerman-Tam, Diana, Arthur Grimes, and Nicholas Watson. 2023. "The economics of free speech: Subjective wellbeing and empowerment of marginalized citizens." *Journal of Economic Behavior & Organization*, 212: 260–274.

von Siemens, Ferdinand A. 2020. "I care what you think: Social image concerns and the strategic revelation of past pro-social behavior." *Journal of the Economic Science Association*, 6: 43–56.

A Appendix (For Online Publication)

A.1 Speaking-up probability is maximized at the expressive norm n

An individual i with private opinion o speaks up if

$$\varepsilon_i \geq \frac{\beta\alpha}{\beta+\alpha}(o-n)^2 - \kappa$$

Define the share of individuals with opinion o who speak up as

$$v^*(o) = \Pr\left(\varepsilon \geq \frac{\beta\alpha}{\beta+\alpha}(o-n)^2 - \kappa\right) = 1 - F_\varepsilon\left(\frac{\beta\alpha}{\beta+\alpha}(o-n)^2 - \kappa\right),$$

where $F_\varepsilon(\cdot)$ is the cumulative distribution function of ε.

Under the assumption that ε is independent of o, and since $F_\varepsilon(\cdot)$ is strictly increasing, willingness to speak up is maximized when the term $(o-n)^2$ is minimized. Since $(o-n)^2$ is uniquely minimized at $o = n$, it follows that

$$\arg\max_o v^*(o) = n.$$

Model extension: expressive utility from ideological extremeness

We extend the baseline model to consider a case in which public expression varies with the strength of (or confidence in) those opinions (Golman, 2023; Le Yaouanq, Schwardmann and van der Weele, 2025).

Suppose that individuals with more ideologically extreme views derive greater intrinsic utility from publicly expressing them, such that: $\varepsilon_i = \zeta_i + g_i$. The idiosyncratic utility from speaking up additionally includes an expressive payoff term g_i which is a function of private opinion o_i such that those with more extreme opinions have a higher payoff from speaking up (perhaps because they hold these views with greater conviction, consistent with the evidence we present on topic importance), and ζ_i is a residual term uncorrelated with o_i. In particular, consider $g_i = g(|o_i - 0.5|)$, where $g'(\cdot) > 0$. The speaking-up condition then becomes:

$$\zeta_i \geq \frac{\beta\alpha}{\beta+\alpha}(o_i-n)^2 - \kappa - g(|o_i - 0.5|),$$

implying

$$v^*(o) = 1 - F_\zeta\left(\frac{\beta\alpha}{\beta+\alpha}(o-n)^2 - \kappa - g(|o-0.5|)\right).$$

The first term inside the cumulative distribution function captures the social cost of deviating from the expressive norm and is minimized at $o = n$. The second term κ is the constant payoff from speaking up, while the third term captures expressive utility from ideological extremeness and increases with distance from the center of the opinion space.

As a result, willingness to speak may initially decline as opinions move away from the expressive norm, but increase again among ideological extremists if expressive utility becomes sufficiently strong. This generates a U-shaped relationship between ideological extremeness and willingness to speak: individuals with moderate views may be least willing to speak publicly because they face some social costs from deviating from the norm while deriving relatively little expressive utility from expression.

Equation (4) continues to characterize the location of the expressive norm only if the increase in expressive utility from ideological extremeness does not outweigh the increase in social costs away from the norm. This assumption cannot be verified directly. In our data, topic importance increases toward the ideological extremes, suggesting that expressive utility may contribute to observed willingness to speak particularly at the ends of the ideological distribution, potentially obscuring the expressive norm peak. Accordingly, we interpret the observed peak in willingness to speak as suggestive rather than definitive evidence regarding the location of the expressive norm. However, we provide complementary evidence for a progressive expressive norm that does not rely on the location of the willingness peak. Specifically, participants believe that public opinion lies to the left of private opinion (see Table 2), and reports of self-censoring are minimized at the extreme left (see Figure A3).

A.2 Weak and Strong Prime treatments

In Experiment 2 (Figure 3), we varied the perceived strength of the prime text. As specified in our pre-registration, this design was motivated by a preliminary finding suggesting that some participants responded to the prime by becoming more willing to "speak up", a pattern consistent with a backlash effect. We therefore sought to examine whether narratives surrounding "cancel culture" might provide a rationale for this response. To do so, we split our prime intervention into a "strong" and a "weak" version. The StrongPrime treatment was identical to the original intervention.

The WeakPrime treatment used the same text while removing bold font, the image, and the text "in a phenomenon that some people refer to as **"cancel culture"**". We again include an attention check after the text, asking: "To check that you are paying attention, what does the text say [cancel culture/social media backlash] can result in?" Participants select from the following alternatives: losing a job, lower voter turnout, or toppling a famous figure, and they cannot proceed unless they select "losing a job". As before, in the Control treatment, participants were shown a text about University College Dublin (UCD) followed by an attention check.

We found no evidence of backlash in the Strong Prime treatment and the effects of the Strong and Weak primes are not statistically different from each other. We therefore pool both treatments throughout our analyses.

A.3 LLM classification methodology

In turn, we provided ChatGPT (version 4.0) with the sample of responses to the two open text questions, "Why did you decide to let us post one or more of your opinions on social media?" and "Why did you decide not to let us post any of your opinions on social media?", and asked it to create a broad categorization of rationales provided by the participants. We then aggregated the categorization into five main types of reason for each action (post or not post), including "Other", and used the OpenAI API to classify each response using the prompts below:

> In a study we conducted, we asked participants whether they would allow us to post their opinion and name on potentially controversial topics (about gender and race) on social media. If they said yes, we asked them for a reason. Categorize the following reason for saying yes –*text*– into one of these: 1. Indifference to Social Pressure: Lack of concern for public reaction or social backlash, not afraid or ashamed of sharing their opinion. 2. Importance of Issues: Reflects strong beliefs about these social issues (gender and race), regardless of stance. 3. Freedom of Speech: Valuing the right and importance of free speech. 4. Financial Incentives: Rewards and incentives from being part of a study and potential financial gain. 5. Other. In your output, write only the number corresponding to the category, no text.
>
> In a study we conducted [...] If they said no, we asked them for a reason. Categorize the following reason for saying no –*text*– into one of these: 1. Privacy and Security Concerns: Concern about personal information and privacy and a desire for minimal digital presence. 2. Fear of Social or Professional Backlash: Worry about negative consequences from opinion being public or being misrepresented. 3. Disinterest or Lack of Confidence in the Issues: Reflects indifference

> or lack of knowledge about social issues (gender and race). 4. Disengagement from Social Media: Reflects low engagement or activity on social platforms due to personal preference and lack of interest. 5. Other. In your output, write only the number corresponding to the category, no text.

We used both GPT 4.0 and GPT 3.5-turbo to classify each of the statements using these prompts. Examples of participants' classified responses are shown in Appendix Table A22.

B Appendix Tables

Table A1: Summary statistics

	Wave 1					*Wave 2*				
	N	Mean	SD	Min	Max	N	Mean	SD	Min	Max
Age	900	28.48	9.45	18	74	750	36.36	13.43	18	81
Male	900	0.25	0.43	0	1	750	0.49	0.50	0	1
White	900	0.77	0.42	0	1	750	0.72	0.45	0	1
College degree	900	0.65	0.48	0	1	750	0.68	0.47	0	1
Employed	900	0.72	0.45	0	1	750	0.75	0.43	0	1
Risk attitude	900	6.35	1.84	0	10	750	5.72	2.14	0	10
Political leaning	900	4.50	2.89	0	10	750	3.77	2.85	0	10
Active SM user	900	0.80	0.40	0	1	750	0.55	0.50	0	1
	Wave 3					*Wave 4*				
	N	Mean	SD	Min	Max	N	Mean	SD	Min	Max
Age	1502	42.08	13.42	18	85	369	38.74	13.68	18	85
Male	1502	0.49	0.50	0	1	369	0.48	0.50	0	1
White	1502	0.79	0.41	0	1	369	0.70	0.46	0	1
College degree	1502	0.63	0.48	0	1	369	0.56	0.50	0	1
Employed	1502	0.75	0.43	0	1	369	0.73	0.45	0	1
Risk attitude	1502	5.11	2.43	0	10	369	5.14	2.35	0	10
Political leaning	1502	4.47	2.94	0	10	369	4.33	2.08	0	10
Active SM user	1502	0.47	0.50	0	1	369	0.47	0.50	0	1

Notes: *Political leaning* is the response to "In political matters, people talk of 'the left' and 'the right'. How would you place your views on this scale, generally speaking?" (0-10). *Active SM users* indicates the response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour.

Table A2: Correlates of concern about political climate and freedom of speech online

	Wave 1		Wave 2		Wave 3	
	(1)	(2)	(3)	(4)	(5)	(6)
Stated view (s)		1.370***		0.937***		1.175***
		(0.120)		(0.162)		(0.100)
Extremeness $(2(s-0.5))^2$		0.193*		0.108		0.017
		(0.101)		(0.144)		(0.080)
Willing to publish with name	-0.618***	-0.404***	-0.637***	-0.529***	-0.741***	-0.653***
	(0.102)	(0.100)	(0.130)	(0.132)	(0.084)	(0.081)
Age	-0.009	-0.012**	-0.002	-0.007*	-0.000	-0.005*
	(0.006)	(0.006)	(0.004)	(0.004)	(0.003)	(0.003)
Female	-0.423***	-0.314**	0.122	0.229*	0.110	0.161**
	(0.129)	(0.122)	(0.118)	(0.118)	(0.080)	(0.076)
Asian or Pacific Islander	-0.192	-0.088	0.157	0.202	0.355**	0.338**
	(0.263)	(0.259)	(0.190)	(0.187)	(0.154)	(0.149)
Black or African American	-0.116	-0.100	0.273	0.156	0.033	0.084
	(0.165)	(0.160)	(0.175)	(0.175)	(0.173)	(0.167)
Hispanic or Latino	-0.220	-0.207	-0.070	-0.095	0.241	0.194
	(0.198)	(0.178)	(0.212)	(0.209)	(0.199)	(0.194)
Other race	-0.146	-0.149	-0.643	-0.742*	0.370	0.323
	(0.332)	(0.309)	(0.483)	(0.430)	(0.245)	(0.230)
College degree	0.035	0.101	0.317**	0.332***	0.180**	0.215***
	(0.118)	(0.113)	(0.128)	(0.125)	(0.082)	(0.078)
Employed	0.025	-0.004	0.218	0.153	0.086	0.049
	(0.125)	(0.117)	(0.138)	(0.137)	(0.092)	(0.089)
Risk attitude	0.141**	0.085	0.147**	0.124**	0.140***	0.131***
	(0.058)	(0.055)	(0.059)	(0.056)	(0.042)	(0.040)
Active SM users	0.193	0.266**	0.376***	0.423***	0.302***	0.340***
	(0.132)	(0.124)	(0.122)	(0.121)	(0.081)	(0.077)
N	1800	1800	750	750	3004	3004
R-sq	0.050	0.148	0.076	0.118	0.065	0.132

Notes: OLS regressions of the first principal component of responses to end-of-survey concern questions: "How often do you worry that things you post on social media can be misinterpreted?"; "The political climate these days prevents me from saying things I believe because others might find them offensive."; "Are you worried about losing your job or missing out on job opportunities if your political opinions become known?"; "How often do you think social pressure causes people to misrepresent or lie about their political opinions on social media?"; "How often do you think social pressure causes people to refrain or abstain from expressing political opinions on social media?". *Active SM users* indicates the response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour. Fixed effects for survey waves × topic included (not shown). Robust standard errors in parentheses are clustered at the individual level. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A3: Concerns about freedom of speech and experimental interventions (Wave 1)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	-0.228	-0.179	-0.315	-0.312	-0.068	-0.107	-0.145	-0.068
	(0.145)	(0.139)	(0.211)	(0.208)	(0.198)	(0.207)	(0.213)	(0.214)
Awareness	-0.268	-0.293*	-0.388	-0.460*	-0.126	-0.200	0.002	0.014
	(0.185)	(0.177)	(0.247)	(0.253)	(0.240)	(0.245)	(0.279)	(0.282)
Prime x Awareness	0.196	0.176	0.353	0.390	-0.143	-0.054	-0.065	-0.161
	(0.236)	(0.228)	(0.347)	(0.339)	(0.312)	(0.323)	(0.339)	(0.346)
N	1800	1800	586	586	607	607	607	607
R-sq	0.004	0.125	0.009	0.121	0.006	0.042	0.007	0.082
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of the first principal component of responses to end-of-survey concern questions: "How often do you worry that things you post on social media can be misinterpreted?"; "The political climate these days prevents me from saying things I believe because others might find them offensive."; "Are you worried about losing your job or missing out on job opportunities if your political opinions become known?"; "How often do you think social pressure causes people to misrepresent or lie about their political opinions on social media?"; "How often do you think social pressure causes people to refrain or abstain from expressing political opinions on social media?". Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A4: Concerns about freedom of speech and experimental interventions (Wave 3)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	-0.037	-0.003	-0.098	-0.029	-0.047	0.011	0.028	0.038
	(0.119)	(0.114)	(0.197)	(0.182)	(0.156)	(0.154)	(0.162)	(0.159)
HiPeer	-0.104	-0.056	0.026	0.108	-0.065	0.002	-0.144	-0.130
	(0.137)	(0.133)	(0.238)	(0.225)	(0.178)	(0.178)	(0.191)	(0.186)
Prime x HiPeer	0.109	0.057	0.022	-0.089	0.246	0.152	-0.021	0.018
	(0.169)	(0.162)	(0.297)	(0.281)	(0.221)	(0.221)	(0.232)	(0.226)
N	3004	3004	684	684	1001	1001	1319	1319
R-sq	0.000	0.095	0.001	0.101	0.009	0.073	0.006	0.069
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of the first principal component of responses to end-of-survey concern questions: "How often do you worry that things you post on social media can be misinterpreted?"; "The political climate these days prevents me from saying things I believe because others might find them offensive."; "Are you worried about losing your job or missing out on job opportunities if your political opinions become known?"; "How often do you think social pressure causes people to misrepresent or lie about their political opinions on social media?"; "How often do you think social pressure causes people to refrain or abstain from expressing political opinions on social media?". Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A5: Correlates of willingness to publish views online with name

	Wave 1			Wave 3			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Stated view (s)	-0.233***	-0.324***	-0.143***	-0.136***	-0.220***	-0.062*	-0.114***
	(0.037)	(0.047)	(0.052)	(0.030)	(0.038)	(0.035)	(0.030)
Extremeness $(2(s-0.5))^2$	0.155***	0.150***	0.153***	0.114***	0.157***	0.087***	0.099***
	(0.029)	(0.037)	(0.041)	(0.022)	(0.030)	(0.029)	(0.022)
Age	0.004**	0.003*	0.005***	0.001	0.001*	0.000	0.000
	(0.002)	(0.002)	(0.002)	(0.001)	(0.001)	(0.001)	(0.001)
Female	-0.110***	-0.094***	-0.117***	-0.016	-0.024	-0.004	-0.006
	(0.033)	(0.036)	(0.037)	(0.020)	(0.022)	(0.022)	(0.020)
Asian or Pacific Islander	-0.050	-0.113**	0.004	-0.042	-0.036	-0.042	-0.030
	(0.048)	(0.053)	(0.058)	(0.040)	(0.043)	(0.042)	(0.039)
Black or African American	0.119**	0.131**	0.104*	0.014	0.013	0.026	0.004
	(0.050)	(0.055)	(0.054)	(0.041)	(0.044)	(0.047)	(0.040)
Hispanic or Latino	0.015	0.002	0.023	0.036	0.076	-0.004	0.035
	(0.051)	(0.055)	(0.056)	(0.047)	(0.052)	(0.049)	(0.047)
Other race	-0.045	-0.114	0.024	-0.013	-0.032	-0.000	-0.020
	(0.087)	(0.089)	(0.101)	(0.068)	(0.070)	(0.073)	(0.069)
College degree	-0.032	-0.052*	-0.011	-0.063***	-0.055**	-0.072***	-0.060***
	(0.029)	(0.032)	(0.031)	(0.021)	(0.023)	(0.022)	(0.021)
Employed	0.021	0.023	0.017	0.038*	0.046*	0.030	0.023
	(0.029)	(0.032)	(0.032)	(0.022)	(0.024)	(0.024)	(0.022)
Risk attitude	0.082***	0.072***	0.088***	0.075***	0.063***	0.084***	0.070***
	(0.013)	(0.014)	(0.014)	(0.010)	(0.012)	(0.011)	(0.010)
Active SM users	0.016	-0.026	0.049	0.062***	0.043*	0.079***	0.032
	(0.034)	(0.039)	(0.036)	(0.020)	(0.022)	(0.022)	(0.021)
Democrat	0.047	0.018	0.072**	0.001	-0.003	0.012	-0.010
	(0.032)	(0.036)	(0.036)	(0.024)	(0.026)	(0.026)	(0.024)
Republican	-0.004	0.014	-0.016	-0.004	0.009	-0.025	-0.013
	(0.031)	(0.035)	(0.034)	(0.025)	(0.028)	(0.027)	(0.025)
Frequency posting							0.054***
							(0.009)
N	1800	900	900	3004	1502	1502	3004
R-sq	0.129	0.149	0.122	0.068	0.072	0.074	0.088
Topic	Both	Gender	Race	Both	Gender	Race	Both

Notes: OLS regressions of willingness to publish with name (0/1) as outcome. *Stated view (s)* is agreement to Gender or Race statement, coded on a 0-1 scale. *Active SM users* indicates the response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour. *Frequency posting* response to "How often on average do you post personal opinions on social media?" (only asked in Wave 3): 1 never, 2 less than once a month, 3 a few (1-3) times a month, 4 a few (1-6) times a week, 5 at least once a day. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A6: Correlates of willingness to pay for online publication

	Wave 1		Wave 2		Wave 3	
	(1)	(2)	(3)	(4)	(5)	(6)
Stated view (s)		-27.235***		-32.972***		-16.247***
		(5.114)		(5.564)		(3.909)
Extremeness $(2(s-0.5))^2$		13.333***		15.495***		13.333***
		(4.241)		(4.213)		(3.113)
Age	0.242	0.272	-0.205	-0.048	-0.212**	-0.175*
	(0.183)	(0.182)	(0.130)	(0.125)	(0.096)	(0.096)
Female	-10.499***	-13.878***	-1.982	-5.309	-2.330	-2.588
	(3.942)	(3.901)	(3.619)	(3.545)	(2.499)	(2.478)
Asian or Pacific Islander	-3.772	-2.111	-6.749	-2.654	-9.830**	-7.718*
	(7.659)	(7.535)	(6.442)	(6.245)	(4.680)	(4.661)
Black or African American	8.349	8.489	-1.522	5.329	2.536	0.925
	(5.765)	(5.575)	(5.435)	(5.439)	(4.881)	(4.895)
Hispanic or Latino	-2.931	-1.857	-6.568	-1.090	-0.991	-0.263
	(5.799)	(5.800)	(7.614)	(7.465)	(5.558)	(5.500)
Other race	-2.789	-4.127	21.321	25.485*	9.137	7.872
	(10.634)	(10.251)	(15.544)	(14.004)	(7.713)	(7.555)
College degree	-6.780*	-7.206**	-6.393*	-4.640	-9.035***	-8.933***
	(3.544)	(3.521)	(3.817)	(3.711)	(2.548)	(2.529)
Employed	3.089	2.693	7.490*	8.558**	4.016	4.259
	(3.738)	(3.672)	(4.086)	(3.954)	(2.748)	(2.728)
Risk attitude	10.141***	9.972***	8.202***	8.184***	8.406***	8.566***
	(1.612)	(1.577)	(1.695)	(1.659)	(1.248)	(1.248)
Active SM users	5.337	3.926	12.058***	10.955***	7.887***	6.962***
	(4.183)	(4.084)	(3.775)	(3.716)	(2.475)	(2.442)
Democrat	12.099***	7.364*	7.943**	1.192	5.880**	3.503
	(3.843)	(3.849)	(3.959)	(4.000)	(2.928)	(2.947)
Republican	-5.159	1.036	-6.283	-3.003	3.688	5.124*
	(3.823)	(3.901)	(4.975)	(5.056)	(2.935)	(3.085)
N	900	900	750	750	1502	1502
R-sq	0.102	0.139	0.079	0.138	0.067	0.086

Notes: OLS regressions of willingness to pay as outcome. *Stated view (s)* is agreement to Gender or Race statement, coded on a 0-1 scale. *Active SM users* indicates the response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A7: Correlates of willingness to publish views online

	Wave 1				Wave 2				Wave 3			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)	(12)
Stated view (s)	-0.233***	-0.166***	-0.184***	-0.117***	-0.216***	-0.171***	-0.187***	-0.144***	-0.136***	-0.063**	-0.107***	-0.031
	(0.037)	(0.038)	(0.038)	(0.039)	(0.051)	(0.053)	(0.053)	(0.054)	(0.030)	(0.031)	(0.030)	(0.031)
Extremeness $(2(s-0.5))^2$	0.155***	0.143***	0.104***	0.091***	0.174***	0.155***	0.143***	0.125***	0.114***	0.096***	0.064***	0.042*
	(0.029)	(0.029)	(0.031)	(0.031)	(0.039)	(0.038)	(0.042)	(0.041)	(0.022)	(0.022)	(0.023)	(0.023)
Self-censoring		-0.039***		-0.039***		-0.036***		-0.035***		-0.046***		-0.047***
		(0.008)		(0.008)		(0.009)		(0.009)		(0.006)		(0.005)
Issue importance			0.053***	0.053***			0.027*	0.026*			0.041***	0.044***
			(0.011)	(0.011)			(0.015)	(0.015)			(0.007)	(0.007)
N	1800	1800	1800	1800	750	750	750	750	3004	3004	3004	3004
R-sq	0.129	0.149	0.141	0.162	0.085	0.107	0.090	0.111	0.068	0.107	0.079	0.120

Notes: OLS regressions of willingness to publish with name (0/1) as outcome. *Stated view (s)* is agreement to Gender or Race statement, coded on a 0-1 scale. *Self-censoring* response to: "The political climate these days prevents me from saying things I believe because others might find them offensive." (1 Strongly disagree - 7 Strongly agree). *Issue importance* response to “How important is the issue discussed in the statement to you?” (1 Not important at all - 5 Extremely important). Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A8: Conformity in stated views in Q2: Heterogeneity by Q1 stance

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Awareness	-0.030	-0.071**	0.014	0.013	-0.041	-0.133**	-0.048	-0.077
	(0.022)	(0.031)	(0.033)	(0.047)	(0.039)	(0.055)	(0.033)	(0.050)
Prime	-0.026	-0.052**	0.003	0.003	-0.014	-0.082*	-0.059*	-0.076**
	(0.020)	(0.023)	(0.029)	(0.034)	(0.036)	(0.043)	(0.032)	(0.038)
Prime x Awareness		0.084*		0.002		0.204***		0.054
		(0.043)		(0.065)		(0.076)		(0.068)
Mean in control	0.539	0.539	0.289	0.289	0.553	0.553	0.771	0.771
N	900	900	293	293	295	295	312	312
R-sq	0.330	0.333	0.431	0.431	0.203	0.224	0.164	0.166
Topic FE	X	X	X	X	X	X	X	X
Controls	X	X	X	X	X	X	X	X

Notes: OLS regressions of stated view s for the second question (in the survey order), defined as agreement to topic statement (reverse-coded for Gender statement) and scaled to 0-1. We classify respondents based on their stance in the first question, as Progressive if their stated view is 1 or 2 (on the 1–7 scale left-right scale), Conservative if it is 6 or 7, and Moderate otherwise (3, 4, or 5). Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A9: Willingness to pay for online publication and experimental interventions

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
				Wave 1				
Prime	2.502	4.032	13.234*	13.984**	-0.575	-0.872	-6.339	-2.068
	(3.655)	(3.566)	(7.004)	(6.994)	(5.899)	(5.915)	(6.058)	(5.945)
Awareness	0.345	1.400	6.673	4.716	-5.859	-4.273	-0.070	4.159
	(3.673)	(3.514)	(6.668)	(6.591)	(5.890)	(6.150)	(6.195)	(5.887)
N	900	900	293	293	312	312	295	295
R-sq	0.013	0.121	0.014	0.101	0.034	0.142	0.042	0.170
				Wave 2				
Prime	-2.758	-2.789	-9.824	-10.229	1.743	2.614	-2.075	-3.470
	(3.525)	(3.408)	(6.328)	(6.522)	(5.642)	(5.468)	(5.609)	(5.537)
N	750	750	232	232	249	249	269	269
R-sq	0.001	0.088	0.010	0.074	0.000	0.158	0.001	0.075
				Wave 3				
Prime	-6.067**	-4.845*	-8.642	-10.651*	-2.735	-1.417	-7.648*	-6.416
	(2.618)	(2.564)	(5.693)	(5.461)	(4.265)	(4.261)	(3.981)	(3.975)
HiPeer	4.173*	2.892	2.210	1.033	1.119	-0.566	7.030*	4.058
	(2.432)	(2.366)	(5.355)	(5.197)	(4.009)	(3.925)	(3.657)	(3.695)
N	1502	1502	341	341	503	503	658	658
R-sq	0.006	0.079	0.024	0.166	0.001	0.083	0.011	0.066
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to pay for publication. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A10: Willingness to publish and experimental interventions (Wave 1)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	0.049	0.053	0.119	0.134	0.013	0.008	0.000	0.015
	(0.042)	(0.039)	(0.078)	(0.082)	(0.056)	(0.049)	(0.062)	(0.058)
Awareness	0.037	0.056	0.051	0.080	0.014	0.002	0.028	0.077
	(0.047)	(0.045)	(0.083)	(0.083)	(0.059)	(0.059)	(0.072)	(0.067)
Mean in control	0.213	0.213	0.303	0.303	0.151	0.151	0.204	0.204
N	994	994	319	319	349	349	326	326
R-sq	0.002	0.124	0.010	0.081	0.002	0.179	0.010	0.199
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to publish with name (0/1). The sample excludes treatment branches in which LR-scale (s) is endogenous, thus keeping only Treatments 2, 3, and 5 in Wave 1 (see Figure 2). Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A11: Willingness to publish and experimental interventions (Wave 3), controling for posting frequency

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	-0.051**	-0.025	-0.022	-0.040	-0.058*	-0.034	-0.071**	-0.027
	(0.021)	(0.028)	(0.044)	(0.060)	(0.030)	(0.044)	(0.030)	(0.038)
HiPeer	0.041**	0.076**	0.045	0.019	-0.002	0.029	0.063**	0.128**
	(0.019)	(0.036)	(0.041)	(0.073)	(0.027)	(0.052)	(0.027)	(0.051)
Prime x HiPeer		-0.051		0.038		-0.047		-0.095
		(0.043)		(0.089)		(0.061)		(0.061)
Mean in control	0.230	0.230	0.333	0.333	0.205	0.205	0.202	0.202
N	3004	3004	684	684	1001	1001	1319	1319
R-sq	0.084	0.085	0.126	0.126	0.067	0.067	0.085	0.088
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to publish with name (0/1). Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, topic FE, and frequency of posting opinions online ("How often on average do you post personal opinions on social media?"). We classify respondents as Progressive if their stated view is 1 or 2 (on the 1–7 scale left-right scale), Conservative if it is 6 or 7, and Moderate otherwise (3, 4, or 5). Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A12: Willingness to publish and experimental interventions

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
					Wave 2			
Prime	-0.048	-0.048	-0.091	-0.101	-0.011	-0.008	-0.047	-0.055
	(0.035)	(0.034)	(0.068)	(0.070)	(0.049)	(0.049)	(0.057)	(0.058)
N	634	634	207	207	216	216	211	211
R-sq	0.003	0.058	0.009	0.072	0.000	0.136	0.003	0.088
					Wave 3			
Prime	-0.058**	-0.045*	-0.013	-0.019	-0.060*	-0.050	-0.083**	-0.070**
	(0.023)	(0.023)	(0.048)	(0.047)	(0.032)	(0.032)	(0.034)	(0.034)
HiPeer	0.060***	0.049**	0.070	0.043	0.005	0.003	0.099***	0.078***
	(0.022)	(0.021)	(0.046)	(0.045)	(0.030)	(0.030)	(0.030)	(0.030)
N	2652	2652	613	613	898	898	1141	1141
R-sq	0.009	0.061	0.023	0.098	0.006	0.056	0.024	0.066
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to publish with name (0/1), *excluding those who believed it was extremely unlikely that posts would be published (only asked in Waves 2-3).* Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A13: Willingness to pay for online publication and experimental interventions (Wave 1)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	8.160*	10.387**	23.388**	27.567***	-0.357	-0.388	1.074	5.227
	(4.613)	(4.495)	(9.470)	(9.641)	(7.637)	(7.281)	(7.252)	(6.978)
Awareness	9.199	11.405*	20.853*	23.686**	-5.530	-3.519	13.943	17.875*
	(5.868)	(5.843)	(11.177)	(11.237)	(9.305)	(9.420)	(10.496)	(10.554)
Prime x Awareness	-14.417*	-16.290**	-21.861	-29.393**	-0.572	-1.306	-21.981*	-21.432*
	(7.509)	(7.264)	(13.887)	(13.623)	(12.031)	(12.095)	(12.950)	(12.717)
Mean in control	-62.654	-62.654	-62.667	-62.667	-61.367	-61.367	-63.815	-63.815
N	900	900	293	293	312	312	295	295
R-sq	0.017	0.126	0.023	0.116	0.034	0.142	0.052	0.179
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to pay for publication. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A14: Willingness to pay for online publication and experimental interventions (Wave 3)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	-4.794	-2.859	-8.101	-6.950	-8.324	-8.963	-1.054	1.275
	(3.605)	(3.535)	(8.400)	(7.761)	(5.981)	(6.097)	(5.256)	(5.209)
HiPeer	5.930	5.607	2.928	5.976	-6.221	-10.444	16.735**	15.269**
	(4.349)	(4.285)	(9.402)	(9.259)	(6.980)	(7.020)	(6.721)	(6.755)
Prime x HiPeer	-2.611	-4.050	-1.064	-7.363	11.231	14.909*	-14.094*	-16.384**
	(5.247)	(5.166)	(11.409)	(11.120)	(8.536)	(8.552)	(8.000)	(7.983)
Mean in control	-61.944	-61.944	-47.566	-47.566	-63.897	-63.897	-67.232	-67.232
N	1502	1502	341	341	503	503	658	658
R-sq	0.006	0.080	0.024	0.167	0.005	0.089	0.016	0.073
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to pay for publication. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A15: Willingness to publish and experimental interventions (Wave 3)

	All		Prog		Mod		Cons	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Prime	-0.052	-0.028	-0.056	-0.036	-0.050	-0.042	-0.058	-0.030
	(0.032)	(0.032)	(0.067)	(0.066)	(0.046)	(0.048)	(0.043)	(0.044)
HiPeer	0.068*	0.072*	0.014	0.021	0.018	0.014	0.133**	0.131**
	(0.039)	(0.039)	(0.078)	(0.078)	(0.055)	(0.056)	(0.057)	(0.057)
Prime x HiPeer	-0.012	-0.034	0.084	0.034	-0.020	-0.016	-0.050	-0.079
	(0.047)	(0.047)	(0.096)	(0.096)	(0.065)	(0.066)	(0.067)	(0.067)
Mean in control	0.240	0.240	0.337	0.337	0.209	0.209	0.215	0.215
N	2652	2652	613	613	898	898	1141	1141
R-sq	0.009	0.061	0.024	0.098	0.006	0.056	0.025	0.067
Topic FE	X	X	X	X	X	X	X	X
Controls		X		X		X		X

Notes: OLS regressions of willingness to publish with name (0/1), *excluding those who believed it was extremely unlikely that posts would be published.* Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A16: Willingness to publish and experimental interventions (Wave 3)

	Prime Treatment				Control (Not Prime)			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
HiPeer	0.024	0.035	-0.016	0.046	0.076^{**}	0.007	0.017	0.144^{***}
	(0.024)	(0.050)	(0.032)	(0.032)	(0.037)	(0.075)	(0.055)	(0.052)
N	2020	468	650	902	984	216	351	417
R-sq	0.073	0.150	0.057	0.068	0.063	0.075	0.077	0.089
Sample	All	Prog	Mod	Cons	All	Prog	Mod	Cons
Topic FE	X	X	X	X	X	X	X	X
Controls	X	X	X	X	X	X	X	X

Notes: OLS regressions of willingness to publish with name (0/1). Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A17: Beliefs about stated views and experimental interventions (Wave 3)

	Belief about views of all participants				Belief about views of those who publish				Perceived gap, $b_{all} - b_{pub}$			
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)	(12)
HiPeer	-0.024	-0.076	0.008	-0.003	0.076^{***}	0.060	0.093^{**}	0.083^{*}	-0.100^{***}	-0.136^{*}	-0.084^{*}	-0.085^{*}
	(0.022)	(0.054)	(0.035)	(0.030)	(0.028)	(0.055)	(0.045)	(0.045)	(0.029)	(0.070)	(0.048)	(0.045)
Prime	0.005	0.010	0.033	-0.002	0.036	0.101^{**}	-0.009	0.048	-0.031	-0.091	0.041	-0.050
	(0.019)	(0.048)	(0.029)	(0.025)	(0.025)	(0.048)	(0.042)	(0.038)	(0.026)	(0.060)	(0.044)	(0.039)
Prime x HiPeer	0.010	0.046	-0.019	-0.021	-0.034	-0.068	-0.020	-0.058	0.044	0.114	0.001	0.038
	(0.028)	(0.064)	(0.043)	(0.038)	(0.034)	(0.068)	(0.056)	(0.054)	(0.036)	(0.083)	(0.062)	(0.053)
N	1502	341	503	658	1502	341	503	658	1502	341	503	658
R-sq	0.053	0.037	0.043	0.038	0.043	0.080	0.040	0.052	0.019	0.045	0.049	0.029
Topic FE	X	X	X	X	X	X	X	X	X	X	X	X
Controls	X	X	X	X	X	X	X	X	X	X	X	X
Sample	All	Prog	Mod	Cons	All	Prog	Mod	Cons	All	Prog	Mod	Cons

Notes: OLS regressions of b_{all}, b_{pub}, and $b_{all} - b_{pub}$ on a Left-Right scale (reverse-coded for Gender statement) and scaled to 0-1. In Experiment 1 (Waves 1 and 2) we elicit beliefs about the *average* view of other participants, in Experiment 2 (Wave 3) we elicit beliefs about the *majority* view of other participants. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A18: Reasons for posting and not posting across left-right scale

	No social pressure		Fear of backlash		Importance of issues		Disinterest in issues	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Stated view (s)	-0.105*	-0.039	0.280***	0.244***	0.022	-0.063	-0.003	0.003
	(0.061)	(0.088)	(0.040)	(0.060)	(0.062)	(0.087)	(0.015)	(0.025)
Extremeness $(2(s-0.5))^2$	-0.069	-0.070	-0.031	-0.010	0.241***	0.243***	-0.038***	-0.035***
	(0.056)	(0.061)	(0.036)	(0.038)	(0.055)	(0.057)	(0.012)	(0.013)
Prime	-0.025	-0.024	0.011	0.022	0.056	0.049	0.006	0.005
	(0.044)	(0.044)	(0.027)	(0.027)	(0.043)	(0.042)	(0.011)	(0.011)
HiPeer	0.071*	0.073*	0.025	0.025	-0.076*	-0.082**	-0.005	-0.006
	(0.043)	(0.042)	(0.025)	(0.024)	(0.042)	(0.041)	(0.010)	(0.010)
N	407	407	1095	1095	407	407	1095	1095
R-sq	0.019	0.099	0.036	0.068	0.057	0.171	0.007	0.018
Controls		X		X		X		X
Sample	Post	Post	NoPost	NoPost	Post	Post	NoPost	NoPost

Notes: OLS regressions of stating a particular reason for posting/not posting, average of DV as categorized by GPT 4.0 and GPT 3.5-turbo using responses to "Why did you decide to let us post one or more of your opinions on social media?" or "Why did you decide not to let us post any of your opinions on social media?". These questions are only asked in Wave 3 to the relevant group (respectively, those who chose to publish for at least one of the two issues, $n = 407$, or those who chose not to publish for both issues, $n = 1095$). Stated view (s) is agreement to Gender or Race statement, coded on a 0-1 scale. For those who posted one of the two statements, we used the response to the relevant statement. For those who posted neither or both of the two statements, we used the average response of the two statements. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square and social media use. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A19: Reasons for posting and not posting across left-right scale

	Freedom of speech		Privacy concerns		Financial incentives		SM disengagement	
	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
Stated view (s)	0.060*	0.023	-0.234***	-0.170***	0.057**	0.098***	-0.025	-0.059
	(0.032)	(0.053)	(0.049)	(0.065)	(0.023)	(0.037)	(0.037)	(0.049)
Extremeness $(2(s-0.5))^2$	-0.066**	-0.042	0.071*	0.051	-0.033	-0.039	-0.009	-0.006
	(0.031)	(0.031)	(0.040)	(0.043)	(0.020)	(0.026)	(0.028)	(0.030)
Prime	-0.033	-0.036	-0.033	-0.033	0.000	0.006	0.014	0.006
	(0.027)	(0.028)	(0.029)	(0.029)	(0.021)	(0.020)	(0.021)	(0.022)
HiPeer	0.015	0.016	-0.039	-0.033	-0.011	-0.007	0.008	0.005
	(0.025)	(0.026)	(0.026)	(0.026)	(0.020)	(0.020)	(0.020)	(0.020)
N	407	407	1095	1095	407	407	1095	1095
R-sq	0.021	0.036	0.024	0.037	0.015	0.062	0.001	0.030
Controls		X		X		X		X
Sample	Post	Post	NoPost	NoPost	Post	Post	NoPost	NoPost

Notes: OLS regressions of stating a particular reason for posting/not posting, average of DV as categorized by GPT 4.0 and GPT 3.5-turbo using responses to "Why did you decide to let us post one or more of your opinions on social media?" or "Why did you decide not to let us post any of your opinions on social media?". These questions are only asked in Wave 3 to the relevant group (respectively, those who chose to publish for at least one of the two issues, $n = 407$, or those who chose not to publish for both issues, $n = 1095$). Stated view (s) is agreement to Gender or Race statement, coded on a 0-1 scale. For those who posted one of the two statements, we used the response to the relevant statement. For those who posted neither or both of the two statements, we used the average response of the two statements. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square and social media use. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A20: Number of statements published and reason given

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
HiPeer	0.088**	0.027	0.094***	0.072*	0.090**	0.091**	0.086**	0.056	0.087**
	(0.040)	(0.031)	(0.036)	(0.037)	(0.039)	(0.037)	(0.039)	(0.036)	(0.039)
No social pressure		1.705***							
		(0.051)							
Importance of issues			1.287***						
			(0.058)						
Freedom of speech				1.922***					
				(0.107)					
Financial incentives					1.109***				
					(0.135)				
Fear of backlash						-0.690***			
						(0.032)			
Disinterest in issues							-0.579***		
							(0.057)		
Privacy concerns								-0.753***	
								(0.032)	
SM disengagement									-0.608***
									(0.033)
N	1502	1502	1502	1502	1502	1502	1502	1502	1502
R-sq	0.068	0.436	0.255	0.187	0.089	0.174	0.078	0.223	0.118
Controls	X	X	X	X	X	X	X	X	X

Notes: OLS regressions of number of statements published. Each of the reasons is the average of DV as categorized by GPT 4.0 and GPT 3.5-turbo using responses to "Why did you decide to let us post one or more of your opinions on social media?" or "Why did you decide not to let us post any of your opinions on social media?". These questions are only asked in Wave 3 to the relevant group (respectively, those who chose to publish for at least one of the two issues, $n = 407$, or those who chose not to publish for both issues, $n = 1095$). Thus, the DV equals 0 for the group not asked. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square and social media use. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A21: Number of statements published and reason given

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
Prime	-0.095**	-0.042	-0.087**	-0.062	-0.091**	-0.073*	-0.092**	-0.092**	-0.084**
	(0.043)	(0.034)	(0.039)	(0.040)	(0.043)	(0.041)	(0.043)	(0.039)	(0.042)
No social pressure		1.704***							
		(0.051)							
Importance of issues			1.283***						
			(0.059)						
Freedom of speech				1.918***					
				(0.106)					
Financial incentives					1.098***				
					(0.133)				
Fear of backlash						-0.685***			
						(0.032)			
Disinterest in issues							-0.579***		
							(0.056)		
Privacy concerns								-0.756***	
								(0.032)	
SM disengagement									-0.604***
									(0.033)
N	1502	1502	1502	1502	1502	1502	1502	1502	1502
R-sq	0.068	0.436	0.254	0.186	0.089	0.172	0.078	0.225	0.118
Controls	X	X	X	X	X	X	X	X	X

Notes: OLS regressions of number of statements published. Each of the reasons is the average of DV as categorized by GPT 4.0 and GPT 3.5-turbo using responses to "Why did you decide to let us post one or more of your opinions on social media?" or "Why did you decide not to let us post any of your opinions on social media?". These questions are only asked in Wave 3 to the relevant group (respectively, those who chose to publish for at least one of the two issues, $n = 407$, or those who chose not to publish for both issues, $n = 1095$). Thus, the DV equals 0 for the group not asked. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square and social media use. Robust standard errors in parentheses are clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A22: Examples of participants' reasons for posting/not posting

Reasons for posting	**Reasons for not posting**
Indifference to social pressure • I'm not too concerned with what people think about me or my opinions • I don't see it as a big deal to post my opinion. • I don't care who knows my opinions	Fear of backlash • it has the potential to cause a lot of trouble for me • It would have severe implications for my job should my superiors saw it and might result in my termination. • I would never post controversial opinions on social media because it could cause backlash that I don't want to deal with
Importance of issues • I feel the issue to me is important to comment on. • This is an opinion I feel strongly about, and I have posted similar statements on social media before. • Because it is an important issue that needs to be addressed and not undermined.	Disinterest in issues • The matter is not important to me. • This is not an issue that I'm passionate about. • i do not want to have this conversation with anyone because i do not have a strong opinion either way
Freedom of speech • Because sometimes it just feels good to let your voice be heard. • i like to be heard • I am free to express my opinion and i think it will go a long way to enlighten the youths	Privacy concerns • These are my personal opinions and mine only, I do not want to share it even if my personal information is not involved. • I did not want to post my name or any info on Twitter. • I do not approve of the dissemination of my personal information on Twitter or any other social media.
Financial incentives • I need money and I don't care what strangers on the internet think about my opinions. • To try to win more lottery tickets to win the money- • because I thought i might be more likely to win the lottery	Social media disengagement • I don't post things on Twitter. I wasn't incentivized enough either. • I dont use social media and prefer to keep it that way. • I never post on social media, and don't want to start now.
Other • I wanted to help with this study as much as possible. • I was just curious • It sounded interesting to do.	Other • I am sure others will post. • I can post my own opinion on my own if I want to • Lottery tickets are not worth the hassle, a different story is 100% sure cash or gift card

Notes: Examples of responses to "Why did you decide to let us post one or more of your opinions on social media?" (left) and "Why did you decide not to let us post any of your opinions on social media?" (right), as classified by GPT 4.0. These questions are only asked in Wave 3 to the relevant group (respectively, those who chose to publish for at least one of the two issues, $n = 407$, or those who chose not to publish for both issues, $n = 1095$).

C Appendix Figures

Figure A1: Public expression and attitudes across topics

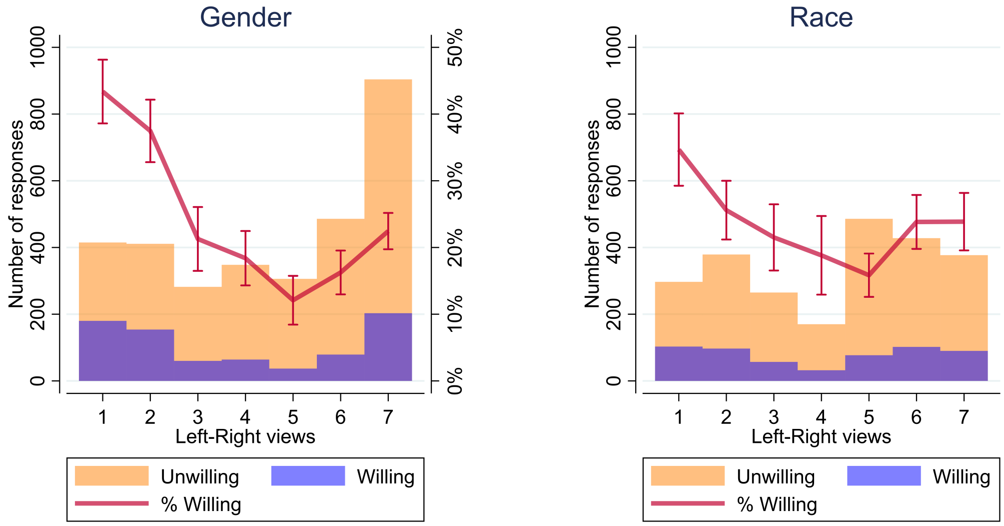

Notes: Agreement to statement in Experiments 1 and 2 (pooled) by willingness to publish and proportion willing to publish. Responses to Gender statement are reverse-coded. Figures show attitudes and public expression separately by topic: Gender (left) and Race (right). Bars represent 95% confidence intervals.

Figure A2: Public expression and attitudes across HiPeer and LoPeer treatments

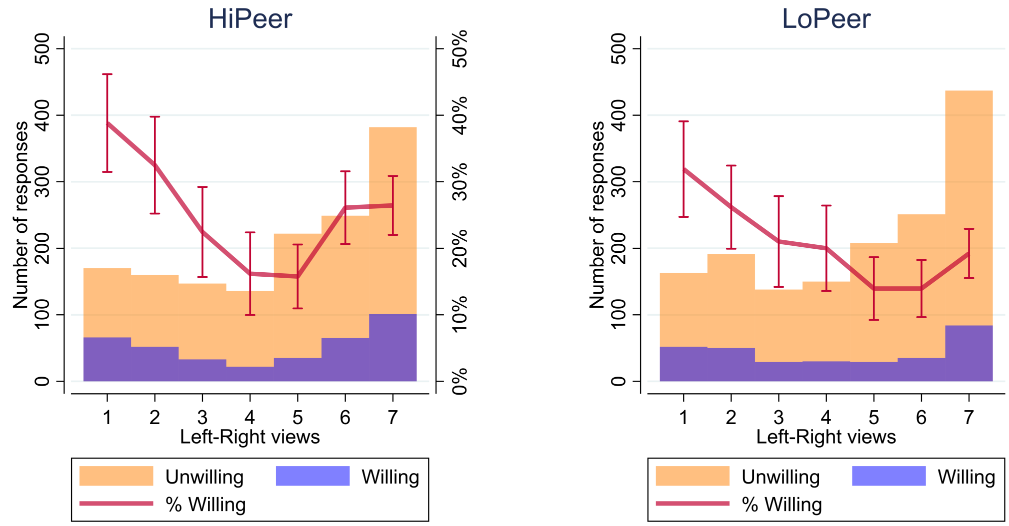

Notes: Agreement to statement in Wave 3 by willingness to publish and proportion willing to publish. Responses to Gender statement are reverse-coded. Figures show attitudes and public expression separately by peer treatment: HiPeer (left) and LoPeer (right). Bars represent 95% confidence intervals.

Figure A3: Issue importance, self-censorship and attitudes across waves

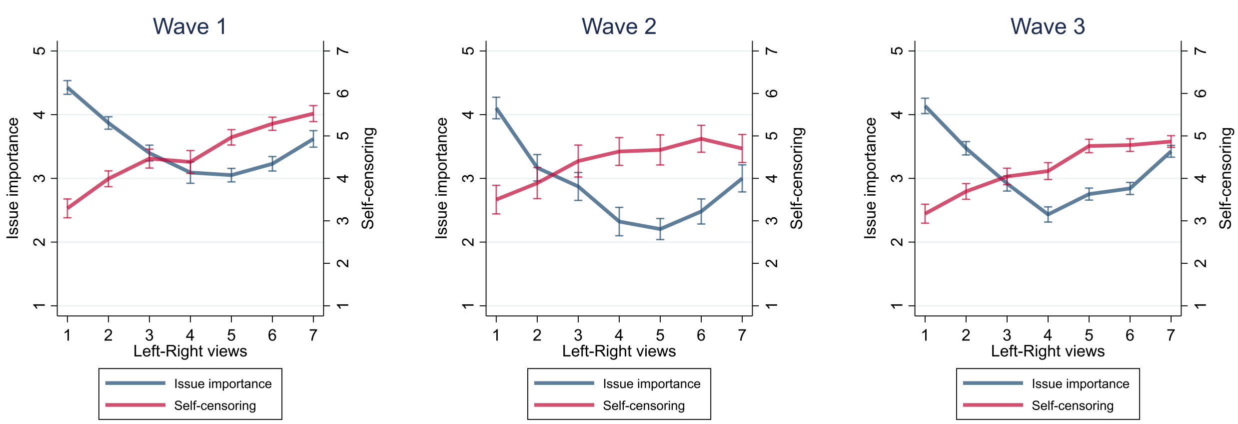

Notes: Issue importance and self-censorship by left-right view. *Issue importance* response to "How important is the issue discussed in the statement to you?" (1 Not important at all - 5 Extremely important). *Self-censoring* response to: "The political climate these days prevents me from saying things I believe because others might find them offensive." (1 Strongly disagree - 7 Strongly agree). *Left-Right view* is agreement to Gender or Race statement. Responses to Gender statement are reverse-coded such that liberal views are on the left. Bars represent 95% confidence intervals.

Figure A4: Heterogeneity of Publication Awareness treatment

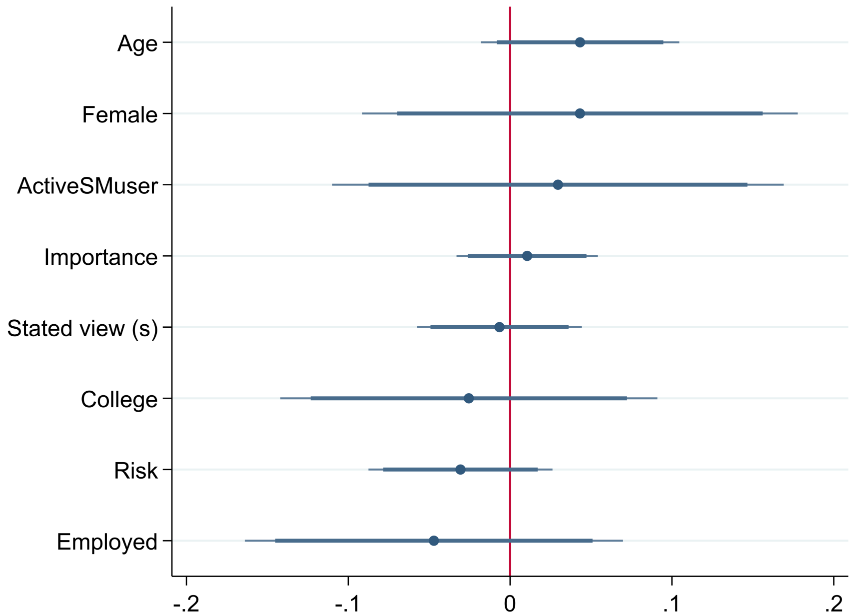

Notes: The figure shows the coefficient θ_3 from the following model:

$$v_{iq} = \theta_0 + \theta_1 Prime_i + \theta_2 Var_i + \theta_3 Awareness_i \times Var_i + \theta_4 Awareness_i + \delta_q + \varepsilon_{iq}$$

where Var_i indicates the dimension of interest in exploring heterogeneous effects. *Stated view (s):* agreement to topic statement (reverse-coded for Gender statement) and scaled to 0-1. *Risk:* response to "Please tell us, in general, how willing or unwilling you are to take risks." (0 Completely unwilling to take risks - 10 Very willing to take risks). *Active SM user:* response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour. *Importance:* response to "How important is the issue discussed in the statement to you?" (1 Not important at all - 5 Extremely important). Except for the dummy variables Female, College, Employed, Active social media use, all other variables are standardised. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors are clustered at the individual level.

Figure A5: Heterogeneity of HiPeer treatment

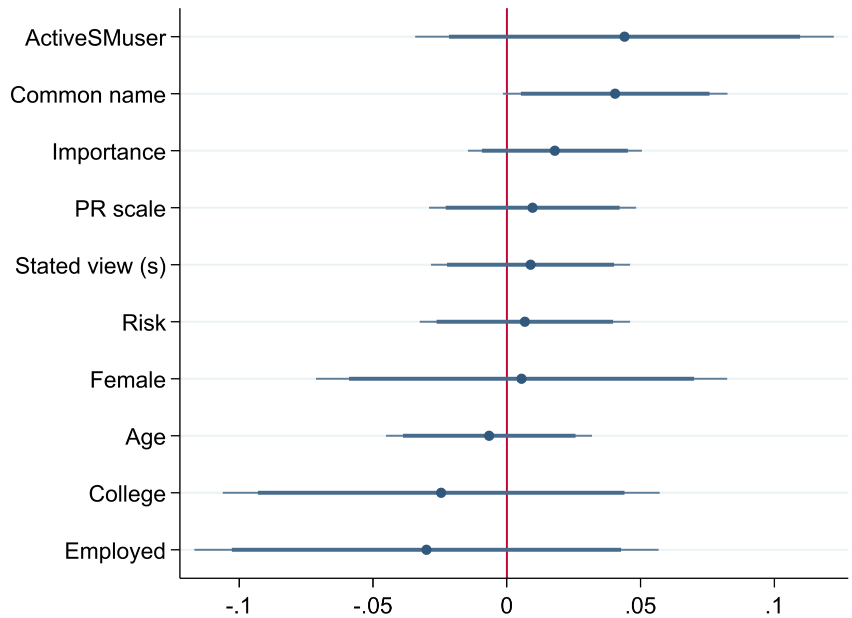

Notes: The figure shows the coefficient θ_3 from the following model:

$$v_{iq} = \theta_0 + \theta_1 Prime_i + \theta_2 Var_i + \theta_3 HiPeer_i \times Var_i + \theta_4 HiPeer_i + \delta_q + \varepsilon_{iq}$$

where Var_i indicates the dimension of interest in exploring heterogeneous effects. *Stated view (s):* agreement to topic statement (reverse-coded for Gender statement) and scaled to 0-1. *Risk:* response to "Please tell us, in general, how willing or unwilling you are to take risks." (0 Completely unwilling to take risks - 10 Very willing to take risks). *Active SM user:* response to "How much time per day do you spend on social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc)" (never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours) and is coded as 1 if the participant states at least 1 hour. *Importance:* response to "How important is the issue discussed in the statement to you?" (1 Not important at all - 5 Extremely important). *Common name:* sum of responses to "The most common first name in the US is James. How common do you think your first name is in the US?" and "The most common last name in the US is Smith. How common do you think your last name is in the US?" (0 Extremely uncommon - 7 Extremely common, Wave 3 only). *PR Scale:* responses to the Hong psychological reactance scale (Hong and Faedda, 1996) (Wave 3 only). Except for the dummy variables Female, College, Employed, Active social media use, all other variables are standardised. Controls include age, gender, race, education, employment, risk attitude, political leaning and its square, social media use, and topic FE. Robust standard errors are clustered at the individual level.

Figure A6: Text analysis of participants' open text-box responses

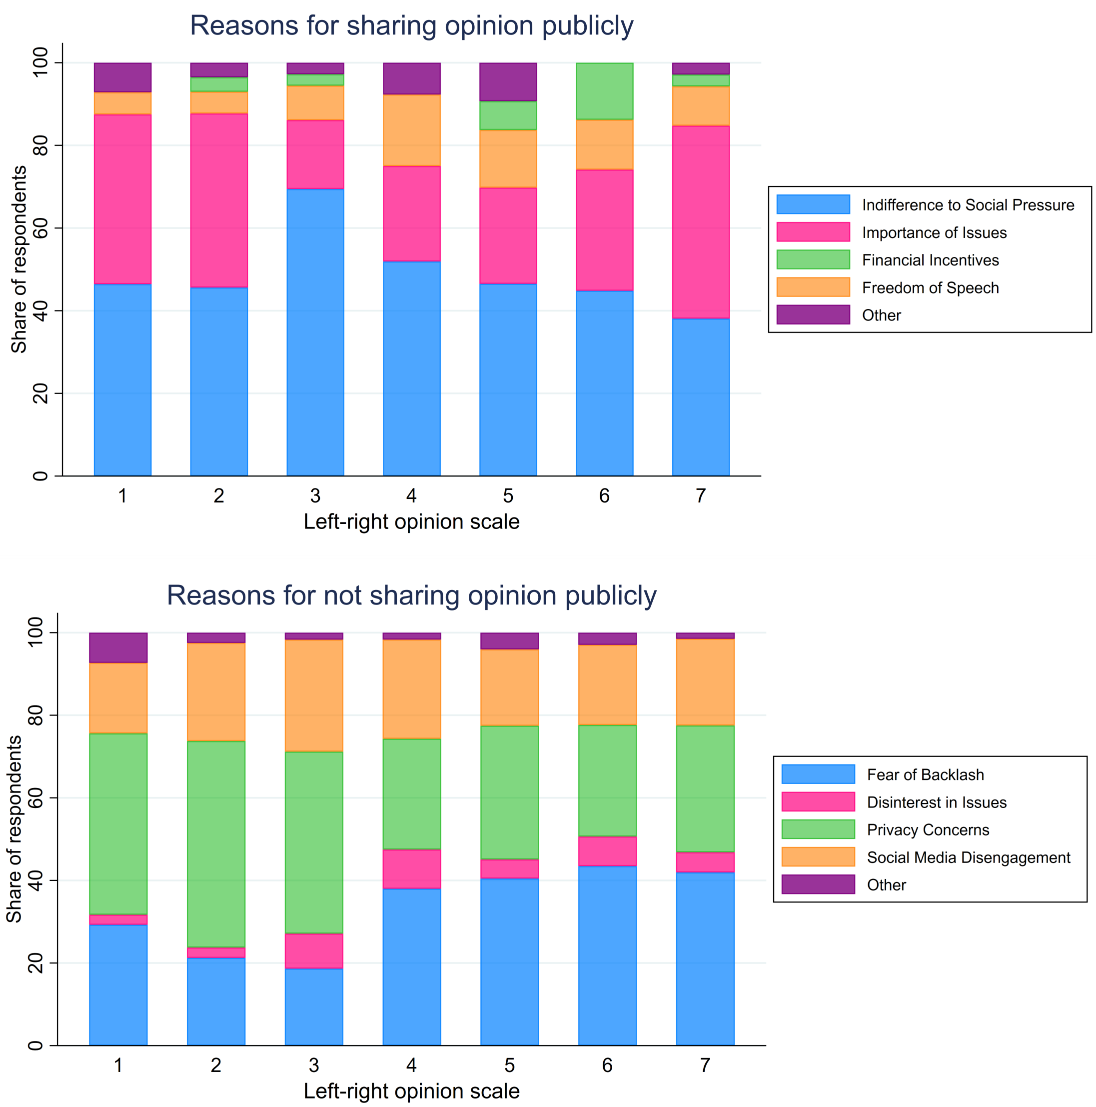

Notes: GPT 4.0 categorization of responses to "Why did you decide to let us post one or more of your opinions on social media?" (top) and "Why did you decide not to let us post any of your opinions on social media?" (bottom). These questions are only asked in Wave 3 to the relevant group (respectively, those who chose to publish for at least one of the two issues, $n = 407$, or those who chose not to publish for both issues, $n = 1095$). Left-Right views are derived from agreement to each of the topic statements, reverse-coded for the Gender topic. For those who posted one of the two statements, we used the response to the relevant statement. For those who posted neither or both of the two statements, we used the average response of the two statements.

Figure A7: Keyword mentions across left-right scale

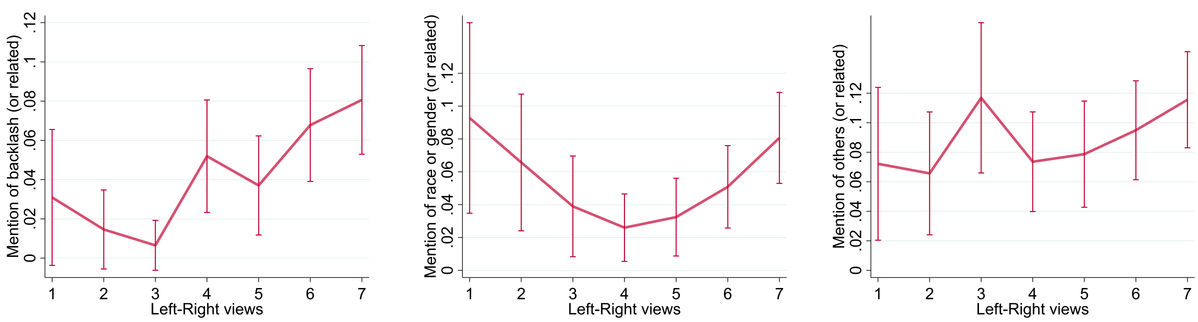

Notes: These figures show the proportion of respondents mentioning certain keywords in their responses to "Why did you decide to let us post one or more of your opinions on social media?" or "Why did you decide not to let us post any of your opinions on social media?". These questions are only asked in Wave 3 to the relevant group (respectively, those who chose to publish for at least one of the two issues, $n = 407$, or those who chose not to publish for both issues, $n = 1095$). The left panel shows the proportion of respondents mentioning “pressure”, “cancel”, “online mob” or “backlash”. The middle panel shows the proportion of respondents mentioning “trans”, “race”, “gender”, “sport” or “lgbt”. The right panel shows the proportion of respondents mentioning “others”, “other people” or “public”. Left-Right views are derived from agreement to each of the topic statements, reverse-coded for the Gender topic. For those who posted one of the two statements, we used the response to the relevant statement. For those who posted neither or both of the two statements, we used the average response of the two statements. Bars represent 95% confidence intervals.

Figure A8: Attitudes in Pilot 1

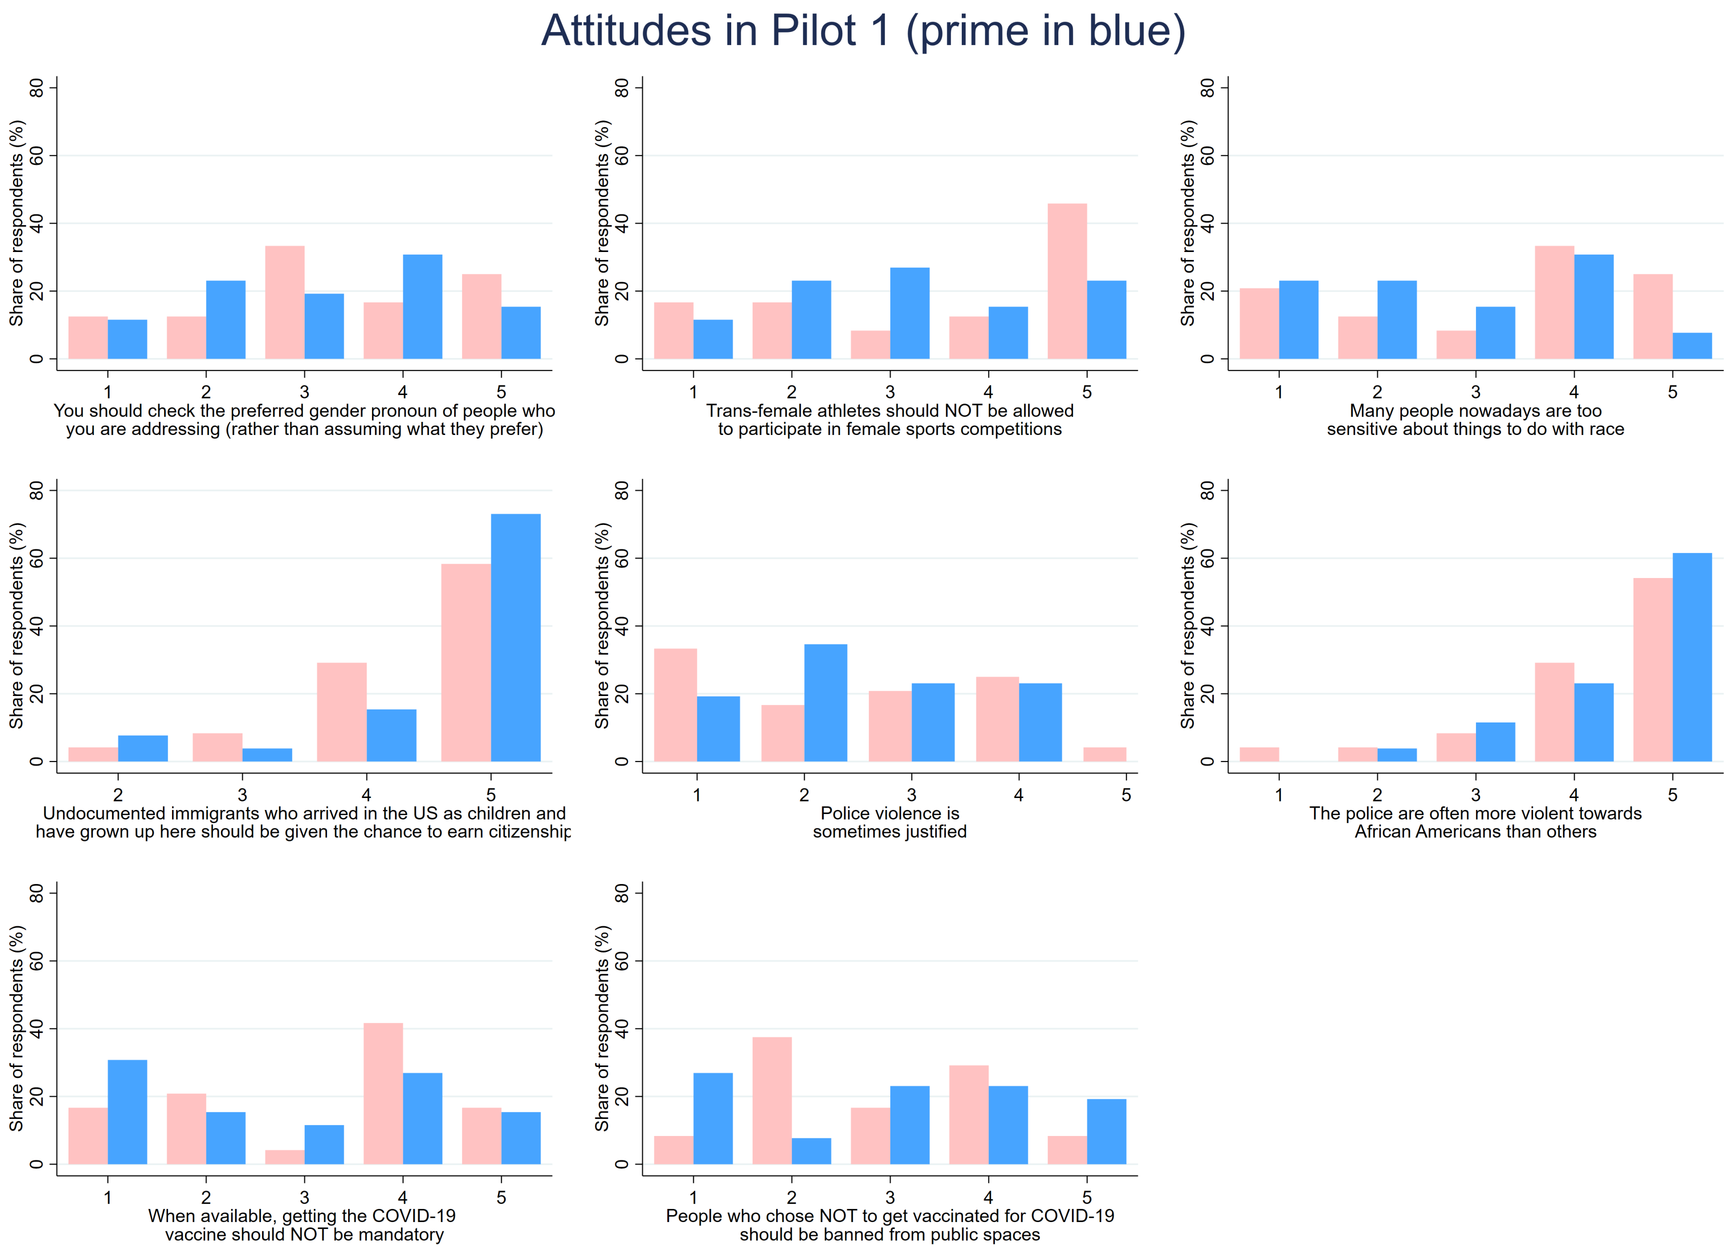

Figure A9: Attitudes in Pilot 2

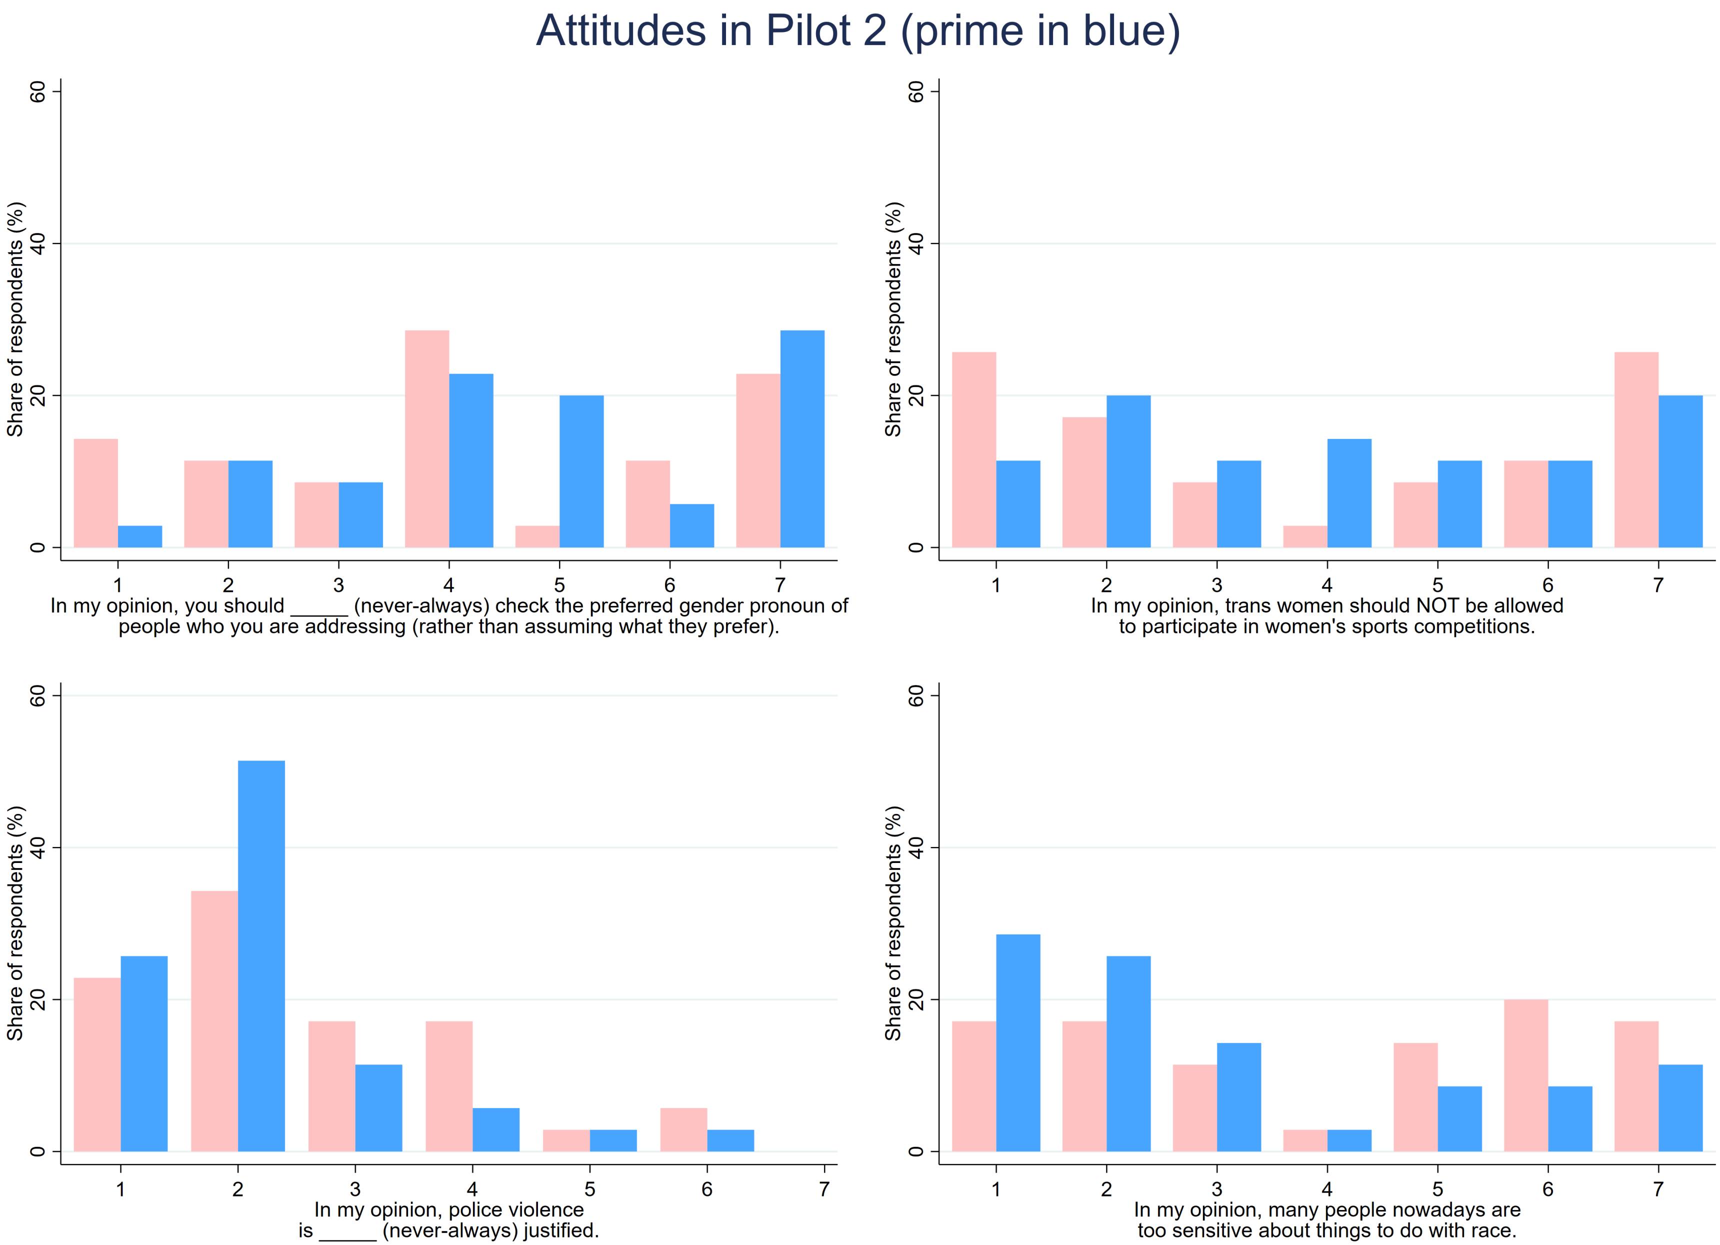

D Full Survey

The full survey for Wave 1 is provided on the next page. Modifications for other waves are detailed in the main text and available with our pre-registration.

[Horizonal lines indicate page break. Unless otherwise specified, all options are presented as radio buttons.]

Introductory Statement
This study is conducted by [redacted]

What is this research about?
This study is part of a research project to study the opinions of Americans.

Why have you been invited to take part?
You have been invited to take part since you meet the research requirement: you are an adult aged over 18 years living in the US.

How will your data be used?
Unless otherwise noted, your data will be analysed and aggregate results will be reported in a future research paper for publication in an academic journal.

What will happen if you decide to take part in this research study?
You will fill out a 10-15 minute survey through Prolific using your desktop computer.

How will your privacy be protected?
Unless otherwise noted, we will collect your Prolific participant ID as is standard procedure, ensuring the data is anonymous.

What are the benefits of taking part in this research study?
Your responses will help researchers better understand the opinions of Americans and how these are formed. You will be paid a participation fee as is standard on Prolific. You will also have the possibility of earning an additional $100 bonus payment through a lottery. You start this survey with 10 tickets and your chance of winning is approximately 1 in 1000.

What are the risks of taking part in this research study?
There are no foreseeable risks to taking part in this study beyond that arising from everyday activities. However, if you have any concern and wish to withdraw at any point, simply close the survey window.

Can you change your mind at any stage and withdraw from the study?
Yes, if you wish to withdraw at any point, simply close the survey window.

How will you find out what happens with this project?
Future updates to the project will be available by contacting the researcher.

Contact details for further information
[redacted]

If you consent to the above information sheet, please select Yes below.
I have read and understood the above and want to participate in this study.
o Yes

o No

Please enter your Prolific ID _____

What is your age (in years)? _____

What is your gender?
- o Man
- o Woman
- o Non-binary/Other ____
- o Prefer not to say

Please specify your ethnicity.
- o White
- o Hispanic or Latino
- o Black or African American
- o Native American or American Indian
- o Asian / Pacific Islander
- o Other ____

In which state do you currently reside? -Dropdown menu containing 50 US states]

What is the highest level of school you have completed or the highest degree you have received?
- o Less than high school degree
- o High school graduate (high school diploma or equivalent including GED)
- o Some college but no degree
- o Associate degree in college (2-year)
- o Bachelor's degree in college (4-year)
- o Master's degree
- o Doctoral degree
- o Professional degree (JD, MD)

Which statement best describes your current employment status?
- o Working (paid employee)
- o Working (self-employed)
- o Not working (temporary layoff from a job)
- o Not working (looking for work)
- o Not working (retired)
- o Not working (disabled)
- o Not working (other) ____
- o Prefer not to answer

Please tell us, in general, how willing or unwilling you are to take risks. [0-10 Likert scale, 0 Completely unwilling to take risks to 10 Very willing to take risks]

In political matters, people talk of ‘the left’ and ‘the right’. How would you place your views on this scale, generally speaking? [0-10 Likert scale, 0 The Left to 10 The Right]

Generally speaking, do you usually think of yourself as a Republican, a Democrat, an Independent, or something else?
o Republican
o Independent
o Democrat
o Other ____
o No preference

How much time per day do you spend...
-On social media (Facebook, Twitter, Instagram, Tik Tok, Snapchat, etc) [never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours]
-Watching, reading or listening to news about politics and current affairs [never/no account, less than 30 minutes, from 30 minutes to 1 hour, from 1 hour to 2 hours, more than 2 hours]

[Control text UNITO]

Please read the following text.

The University of Turin is one of the most ancient and prestigious Italian Universities. Hosting over 79,000 students and with 120 buildings in different areas in Turin and in key places in Piedmont, the University of Turin can be considered as "city-within-a-city", promoting culture and producing research, innovation, training and employment.

Facilities include 22 libraries spread over 32 locations, the Botanic Garden and several University Museums such as "Cesare Lombroso" - Criminal Anthropology Museum and "Luigi Rolando" - Human Anatomy Museum.

To check that you are paying attention, how many museums are named in the text?
o 22
o 2
o 32

[Control text UCD]

Please read the following text.

University College Dublin (commonly referred to as UCD) is a research university in Dublin, Ireland, and a member institution of the National University of Ireland. With 33,284 students, it is Ireland's largest university. Five Nobel Laureates are among UCD's alumni and current and former staff. UCD's main campus is located on a 133-hectare (330-acre) campus at Belfield, four kilometres to the south of the city centre. In 1991, it purchased a second site in Blackrock. This currently houses the Michael Smurfit Graduate Business School.

A report published in May 2015 showed the economic output generated by UCD and its students in Ireland amounted to €1.3 billion annually.

To check that you are paying attention, where does the text say UCD's main campus is located?

- Smurfit
- Belfield
- Blackrock

[Prime text]

Please read the following text.

The public nature of social media has resulted in individuals sometimes experiencing negative consequences as a result of their posts, in a phenomenon that some people refer to as **"cancel culture".**

*"Those most vulnerable to harm tend to be **individuals previously unknown to the public**, like the communications director who was **fired** in 2013 after posting on social media, from her personal account, **an ill-thought-out joke** about Africa, AIDS and her own white privilege ... or the data analyst who was **fired** last spring after posting on social media, after the death of George Floyd in police custody, a study that suggested that riots depressed rather than increased Democratic Party votes."*

These cases highlight the risk of **public backlash from social media.**

To check that you are paying attention, what does the text say cancel culture can result in?
o losing a job
o lower voter turnout
o toppling a famous figure

[OPINION ELICITATION---for a description of the survey logic, please see the experimental design]

You will now be asked to state your opinion on a number of questions.

Please consider the following statement.

People who have been vaccinated against COVID-19 should be allowed to travel without testing and quarantine requirements.

What do you think of the above statement? [1-7 Likert scale, 1 Strongly disagree to 7 Strongly agree]

Please consider the following statement.

In my opinion, trans women should be allowed to participate in women's sports competitions.

What do you think of the above statement? [1-7 Likert scale, 1 Strongly disagree to 7 Strongly agree]

Please consider the following statement.

In my opinion, many people nowadays are too sensitive about things to do with race.

What do you think of the above statement? [1-7 Likert scale, 1 Strongly disagree to 7 Strongly agree]

[PUBLICATION ELICITATION---for a description of the survey logic, please see the experimental design]

Would you be willing to let us post on social media, <u>anonymously</u>, your response to the previous statement:

[Participant 37]
"I [pipe selected choice] that many people nowadays are too sensitive about things to do with race."

[In T4-5 only] Note that if you would like to change your response, you can simply click the Back (left arrow) button below to go back to the previous page.

If you select Yes, **we will create a tweet** containing the above response and post it on a public Twitter page created once data collection is complete. Participant numbers (eg, 37 in the above) are randomly assigned and not linked to your identity in any way.

If you select No, we will NOT create a tweet containing the above.

o Yes

o No

Would you be willing to let us post on social media, <u>together with your name</u>, your response to the previous statement:

[Your name here]

"I [pipe selected choice] that many people nowadays are too sensitive about things to do with race."

[In T4-5 only] Note that if you would like to change your response, you can simply click the Back (left arrow) button below to go back to the previous page.

-**We will create a tweet** containing the above response and may post it on a public Twitter page created once data collection is complete (* see below)
-***We will contact Prolific to request your first and last names.** Note that while in general Prolific does not allow researchers to collect personal information, Prolific does encourage researchers to get in touch in cases such as this, where the study design requires the collection of personal data (see https://researcher-help.prolific.co/hc/en-gb/articles/360015378834-Can-I-ask-Participants-for-their-Personal-Information-Identifiers-).
-The tweet will only contain **a text of your name without any hyperlink**, the public Twitter page will potentially contain the names and opinions of many participants.
-The link to the public Twitter page will be **made available to participants** who contact the researcher to ask for it, but it will not be otherwise advertised. The public Twitter page will be **deleted after 30 days.**

o Yes, I would like to

o No, I'd rather not

[WILLINGNESS TO PAY ELICITATION for subjects who chose "Yes, I would like to" above---for a description of the survey logic, please see the experimental design]

You stated that you would like us to post on social media, <u>together with your name</u>, your response to the previous statement:

[Your name here]

"I [pipe selected choice] that many people nowadays are too sensitive about things to do with race."

In exchange for this post, we want to know if you would be willing to **give up some of your lottery tickets** for the $100 bonus (remember that you start with 10 tickets).

*Would you be willing to give up **all 10 lottery tickets** in exchange for this public post? [This question is repeated with 5 lottery tickets and 1 lottery ticket. If Yes is selected, the subject moves on to the next section.]*

o Yes
o No

If you select Yes,
-**We will contact Prolific to request your first and last names.** Note that while in general Prolific does not allow researchers to collect personal information, Prolific does encourage researchers to get in touch in cases such as this, where the study design requires the collection of personal data (see https://researcher-help.prolific.co/hc/en-gb/articles/360015378834-Can-I-ask-Participants-for-their-Personal-Information-Identifiers-).
-Note, we will only reduce your lottery tickets if we do publish the above text with your name.

[WILLINGNESS TO ACCEPT ELICITATION for subjects who chose "No, I'd rather not" above---for a description of the survey logic, please see the experimental design]

You would rather not let us post on social media, <u>together with your name</u>, your response to the previous statement:

[Your name here]
"I [pipe selected choice] that many people nowadays are too sensitive about things to do with race."

We would now like to ask whether you would be willing to **change your mind** in exchange for a **higher chance of winning the $100 lottery**. Remember that you start with 10 tickets.

*Would you be willing to let us post the above if we give you **1 additional lottery ticket?** [This question is repeated with 5, 25, and 50 lottery tickets. If Yes is selected, the subject moves on to the next section.]*

o Yes
o No

If you select Yes,
-You will get 1 additional ticket in the lottery.
-**We will contact Prolific to request your first and last names.** Note that while in general Prolific does not allow researchers to collect personal information, Prolific does encourage researchers to get in touch in cases such as this, where the study design requires the collection of personal data (see https://researcher-help.prolific.co/hc/en-gb/articles/360015378834-Can-I-ask-Participants-for-their-Personal-Information-Identifiers-).

Please consider the following statement.

In my opinion, trans women should be allowed to participate in women's sports competitions.

How important is the issue discussed in the statement to you? [1-5 Likert scale, 1 Not important at all to 5 Extremely important]

Please consider the following statement.

In my opinion, many people nowadays are too sensitive about things to do with race.

How important is the issue discussed in the statement to you? [1-5 Likert scale, 1 Not important at all to 5 Extremely important]

[NORM ELICITATION---for a description of the survey logic, please see the experimental design]

As earlier mentioned, you have the chance to win an additional bonus of $100 through a lottery.

You will now see **2 questions**. You will earn **5 additional lottery tickets** for each question you answer correctly, in addition to your existing tickets.

Therefore, please consider your answers carefully since each correct answer will increase your chance of winning the $100 bonus.

You will now be asked what you think about the **average** opinion out of other participants in this study.

Here is an example using the COVID-19 question. Suppose that the share of participants who state a particular opinion (between 1 to 7) is as shown in the graph below.

The average opinion is calculated by summing up everyone's opinion and dividing by the total number of participants. In this example, the average opinion is **5 - Somewhat agree.**

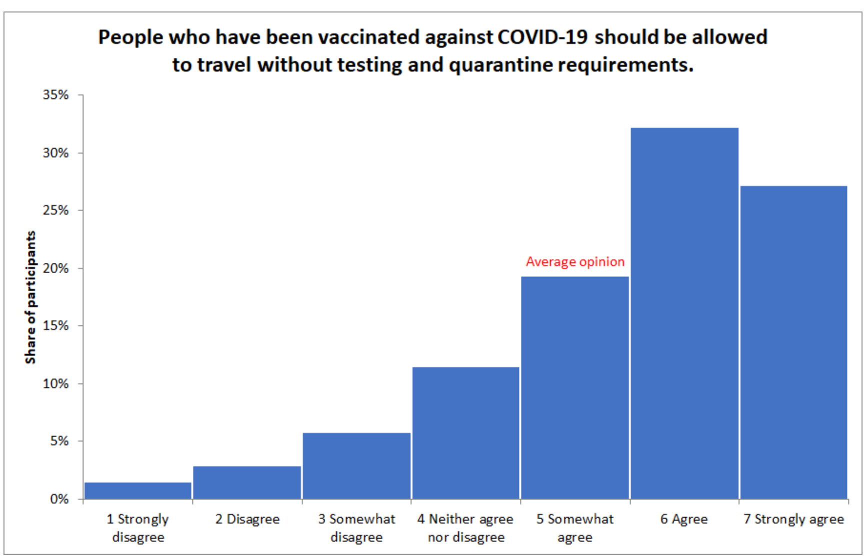

Please consider the following statement.

In my opinion, many people nowadays are too sensitive about things to do with race.

Remember, you will earn 5 additional lottery tickets for each correct answer, so please consider your answers carefully.

Considering ALL participants (in this US-based survey), what do you think the **average** opinion is? [1-7 Likert scale, 1 Strongly disagree to 7 Strongly agree]

Considering those participants (in this US-based survey) who stated that they WOULD be willing to let us post their opinion, together with their name, on social media (without any additional payment), what do you think the **average** opinion is? [1-7 Likert scale, 1 Strongly disagree to 7 Strongly agree]

How often do you worry that things you post on social media can be misinterpreted? [1-7 Likert scale, 1 Never to 7 Always]

The political climate these days prevents me from saying things I believe because others might find them offensive. [1-7 Likert scale, 1 Strongly disagree to 7 Strongly agree]

Are you worried about losing your job or missing out on job opportunities if your political opinions become known? [Not at all worried, Not very worried, Worried a little, Worried a lot]

How often do you think social pressure causes people to **misrepresent or lie about** their political opinions on social media? [1-7 Likert scale, 1 Never to 7 Always]

How often do you think social pressure causes people to **refrain or abstain from expressing** political opinions on social media? [1-7 Likert scale, 1 Never to 7 Always]

Thank you for participating in our study.

This study aims to investigate the impact of cancel culture on self-expression. We are interested in how willing you would be to let us post your opinion on social media.

You were shown some of the following three texts:

- The text about UCD was modified from https://en.wikipedia.org/wiki/University_College_Dublin and serves as a filler.
- The text about UNITO was modified from https://en.unito.it/about-unito/unito-glance and serves as a filler.
- The text about cancel culture was modified from https://www.nytimes.com/2020/12/03/t-magazine/cancel-culture-history.html

As data collection is ongoing, we would like to ask you not to talk about this study with others for now.

If you win the bonus payment, it will be paid through Prolific in the next few weeks.

Regarding the publication of your opinion on social media:

- We will create a public Twitter page for the study.
- We will create an anonymous tweet for each participant's opinion that they are willing to publish.
- Previous requests to Prolific asking for participant's names in a similar study design have been turned down; so we do not anticipate that we will publish your opinion with your name, even if you stated that you would like us to do this. [For subjects who were willing to accept extra tickets for publishing:] Regardless, if you stated that you were willing to publish the opinion with your name in exchange for lottery tickets, you will still get these additional lottery tickets.

If you have any questions about the study, please feel free to contact [redacted]